MARY ROSE LEGASPI DEVELOPMENT PROJECT

MR. MONIR (SUPERVISOR)

BUILDING SCALABLE SOLUTION FOR PREDICTING HEART DISEASE USING APACHE SPARK MLLIB IN STANDALONE CLUSTER MODE WITH MONGODB DATABASE

BACKGROUND AND CONTEXT

WHAT IS SPARK?

What is Apache Spark?

- Apache open sourced project originally developed at AMPLab (UC Berkeley)
- Unified general data processing engine that operates varied data workloads and platforms
- Built on top of Hadoop Map Reduce and it extends the MapReduce model to efficiently use more types of computations

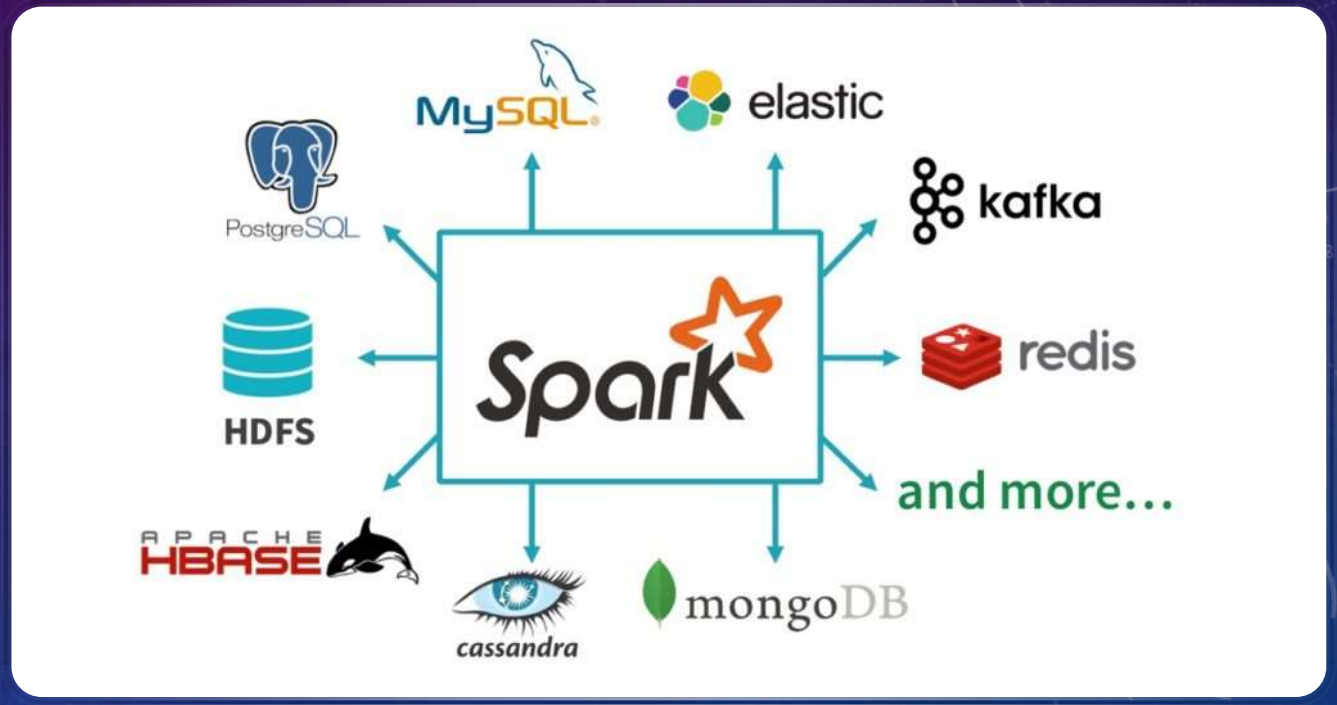

Spark features

100x faster than for large scale data processing

Simple programming layer provides powerful caching and disk persistence capabilities

Can be programmed in Scala, Java, Python, and R.

Can be deployed through Mesos, Yarn, EC2 or Sparks standalone cluster manager

Apache Spark Domain scenarios

Spark Components

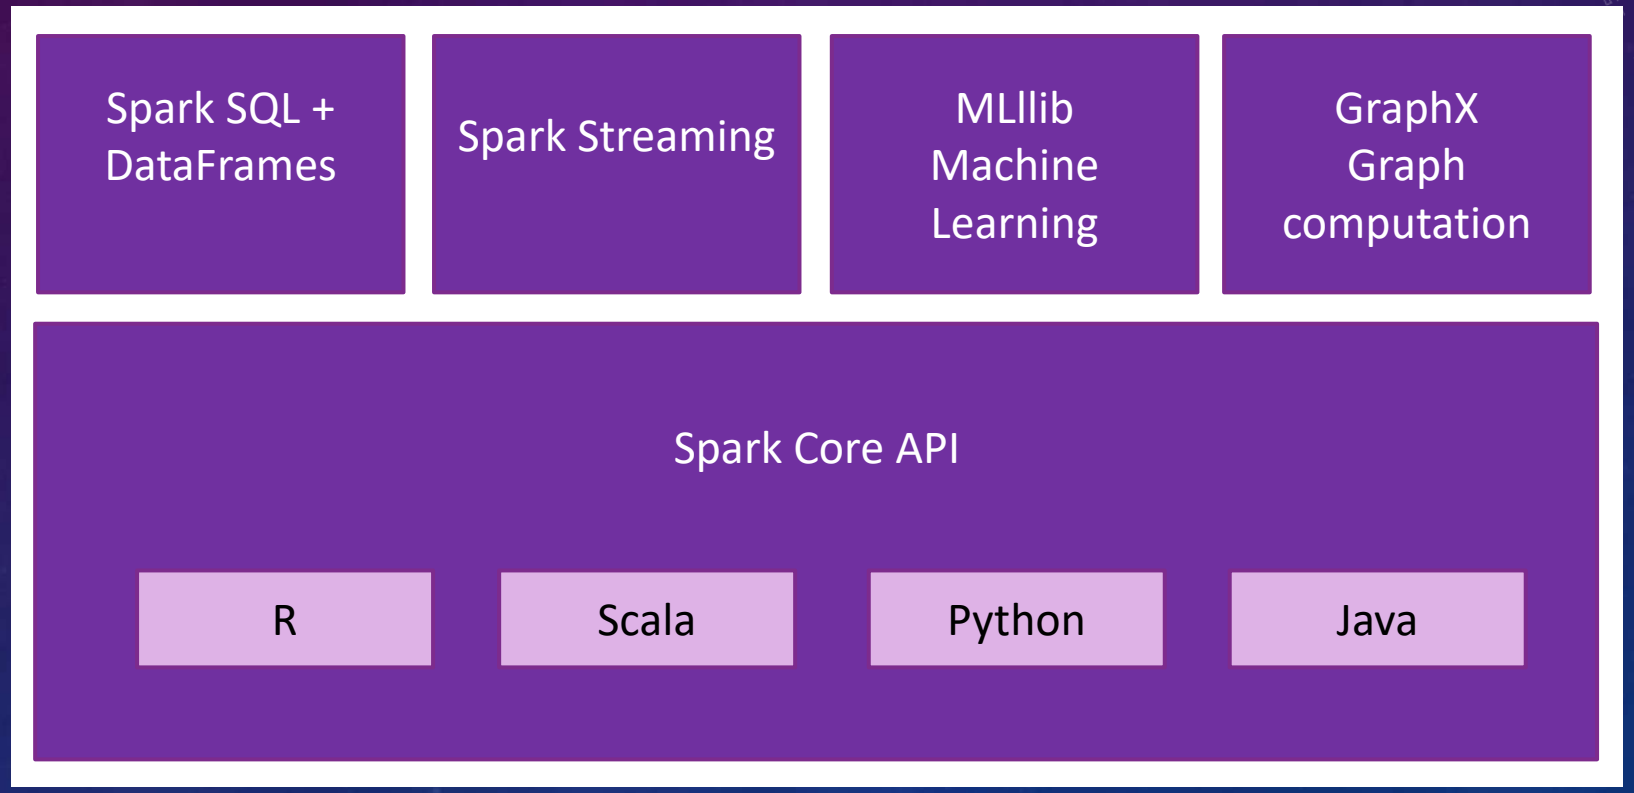

Spark Components

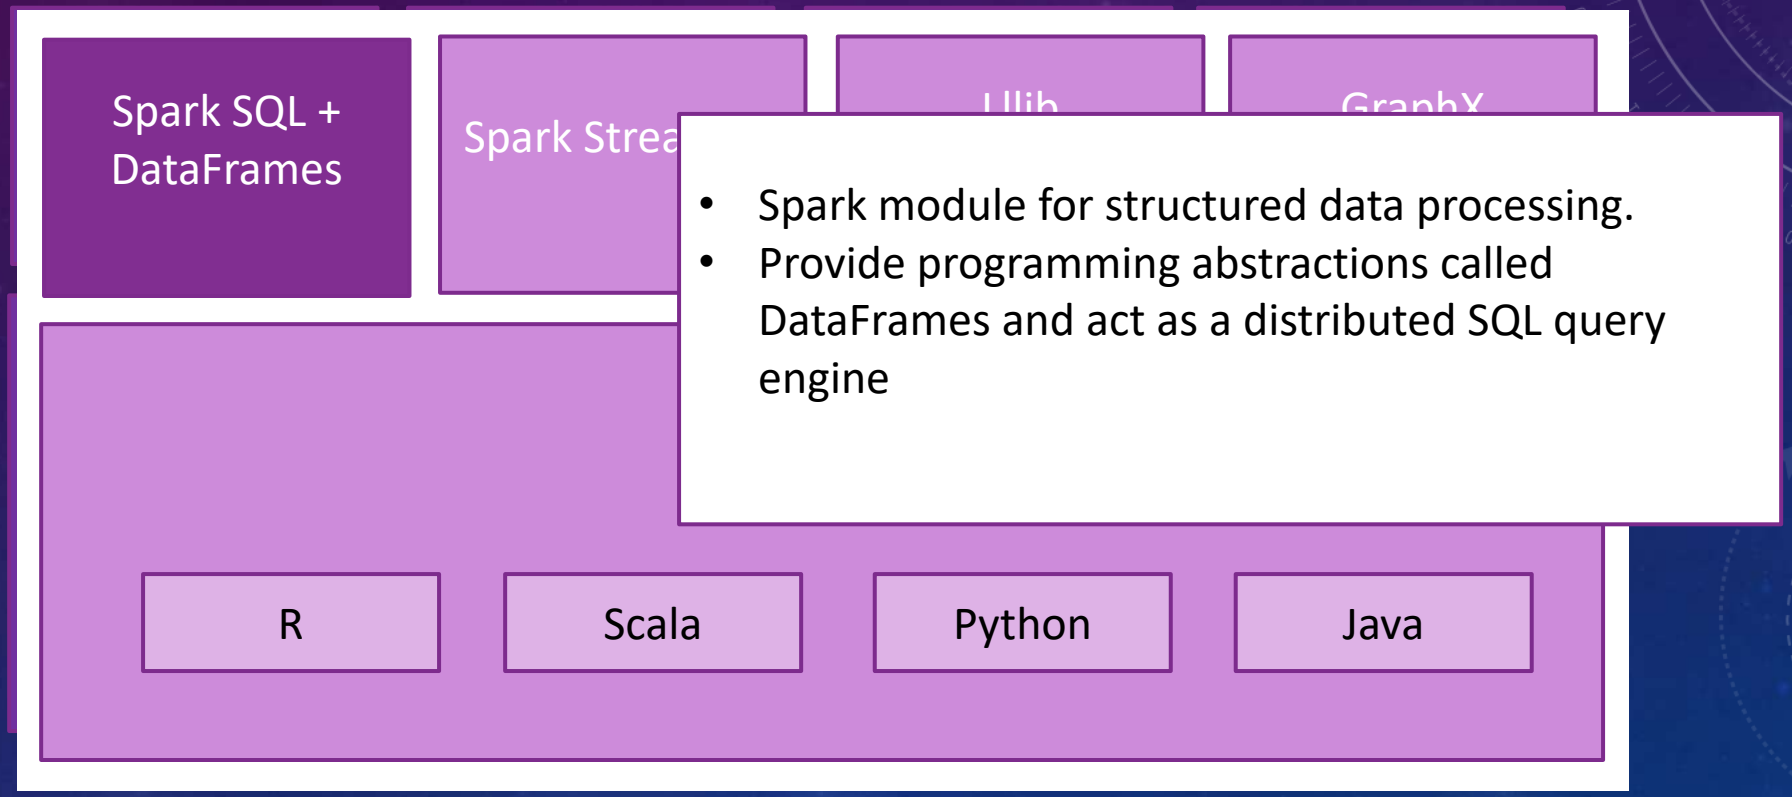

- Spark module for structured data processing.
- Provide programming abstractions called DataFrames and act as a distributed SQL query engine

Spark Components

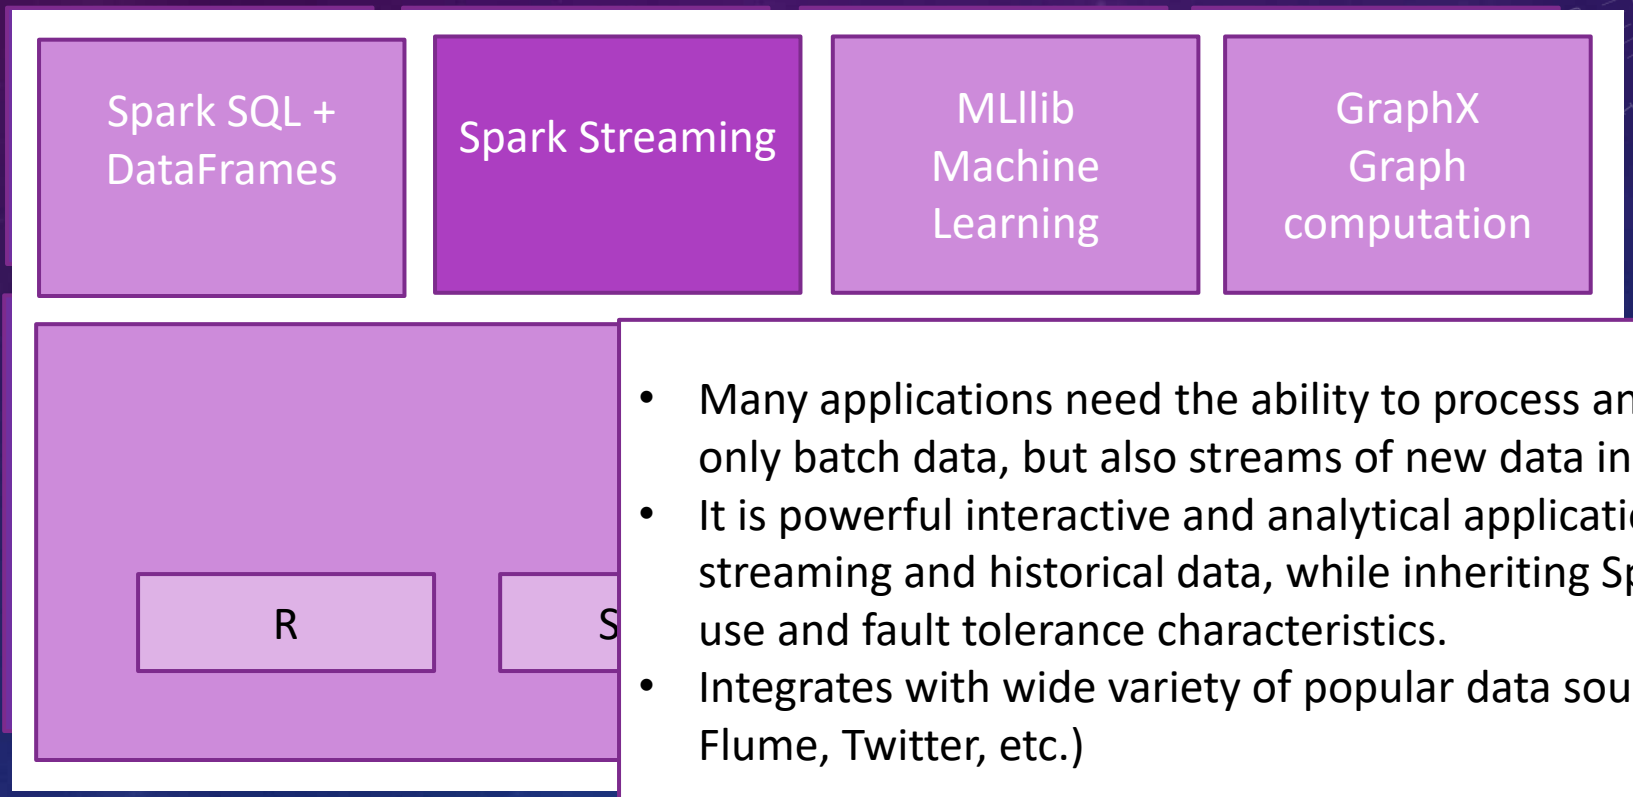

- Many applications need the ability to process and analyse not only batch data, but also streams of new data in real-time
- It is powerful interactive and analytical applications across both streaming and historical data, while inheriting Spark's ease of use and fault tolerance characteristics.
- Integrates with wide variety of popular data sources (HDFS, Flume, Twitter, etc.)

Spark Components

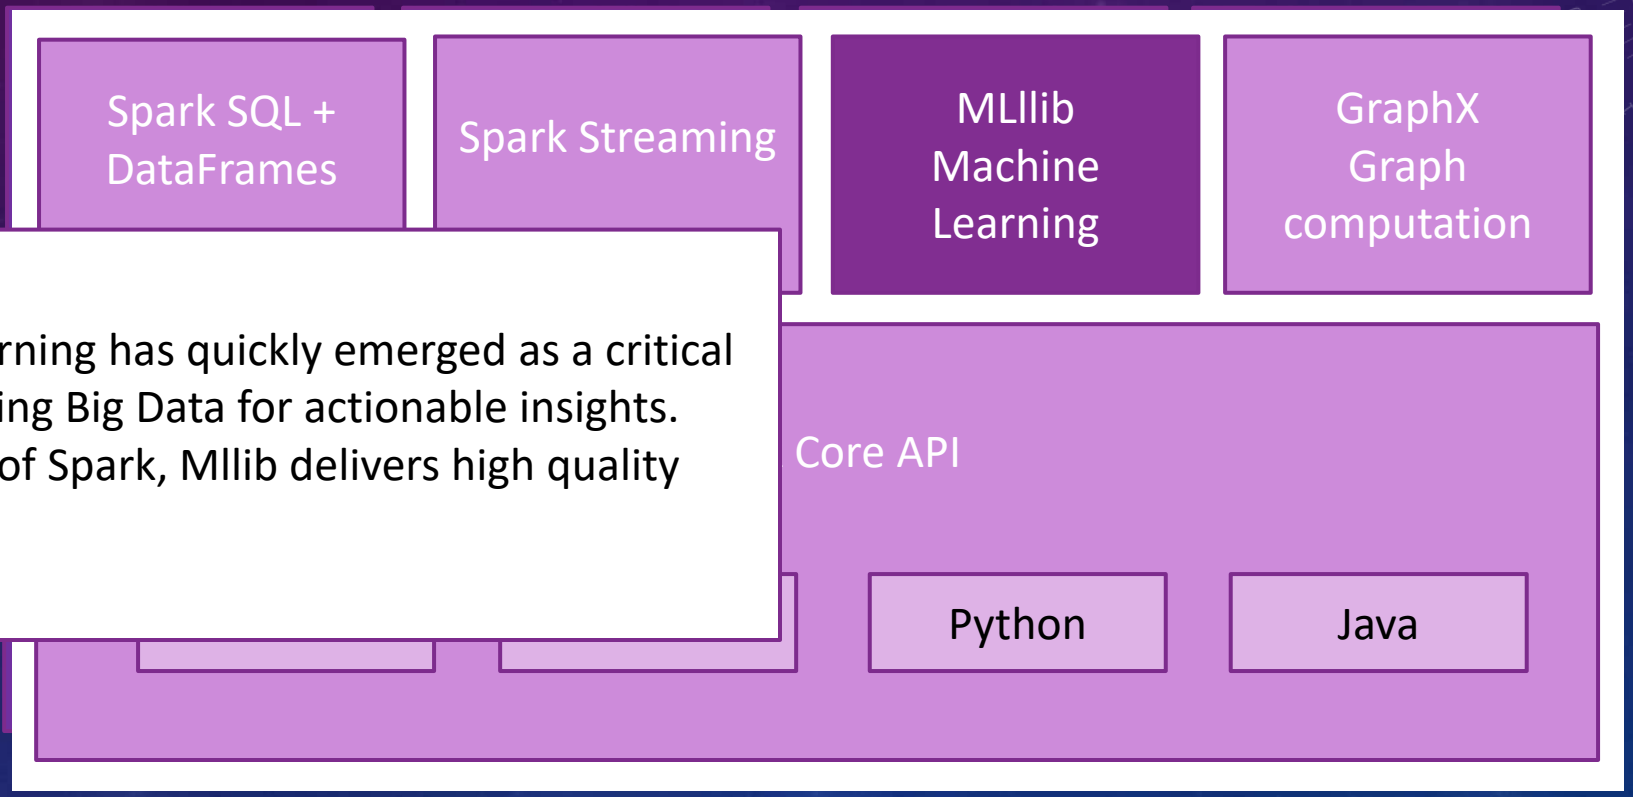

- Machine learning has quickly emerged as a critical piece in mining Big Data for actionable insights. Built on top of Spark, Mllib delivers high quality algorithms.

Spark Components

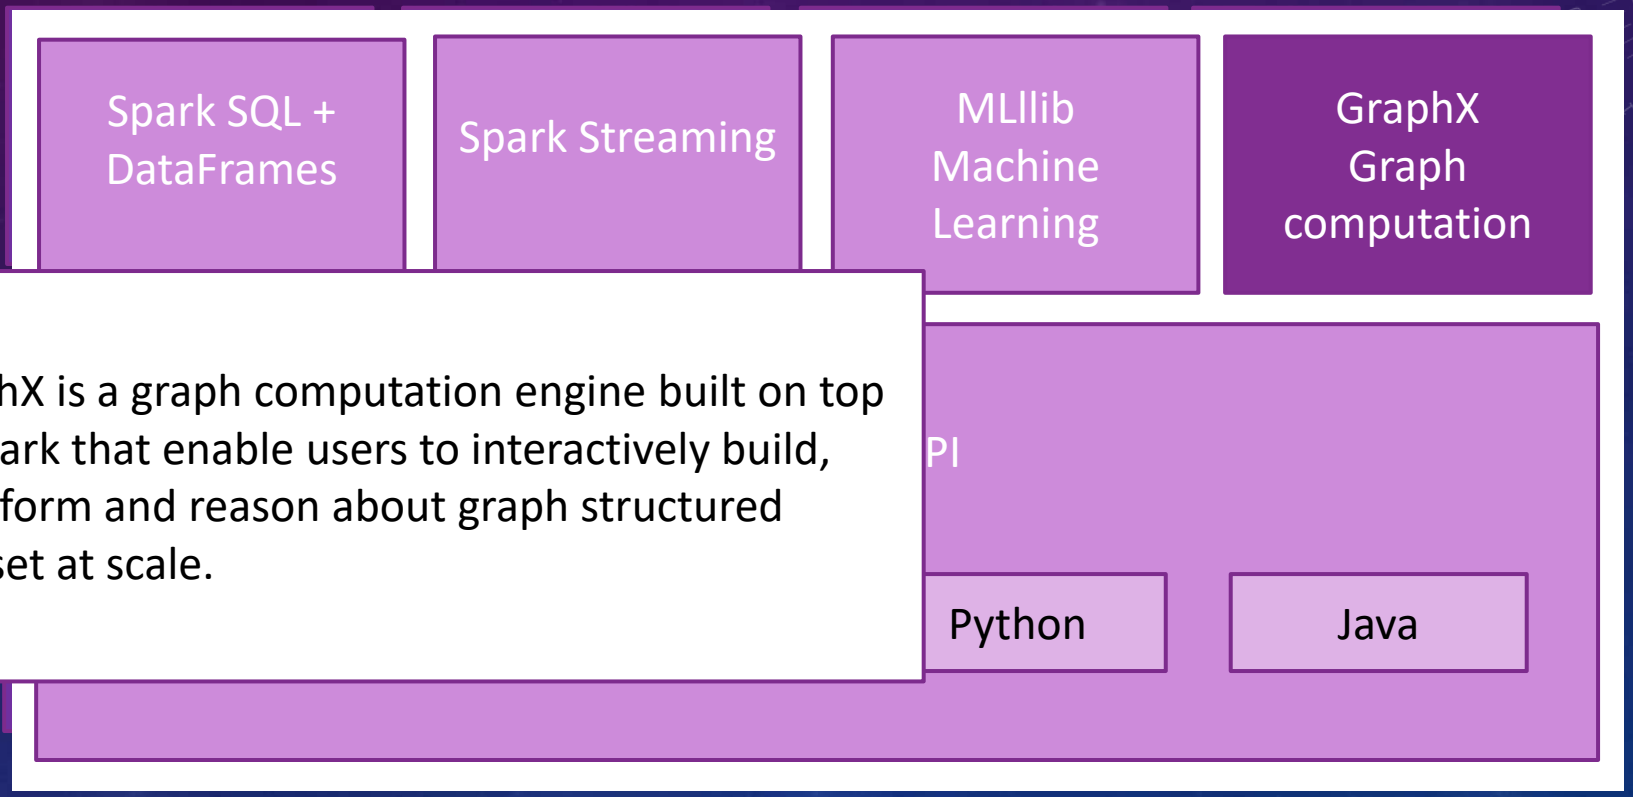

- GraphX is a graph computation engine built on top of Spark that enable users to interactively build, transform and reason about graph structured dataset at scale.

Spark Components

- Spark core is the underlying general execution engine for the Spark platform that all other functionality is built on top of.
- Provides in-memory computing capabilities to deliver speed, a generalized execution model to support a wide variety of applications, and Java, Python, Scala and R API for ease of development

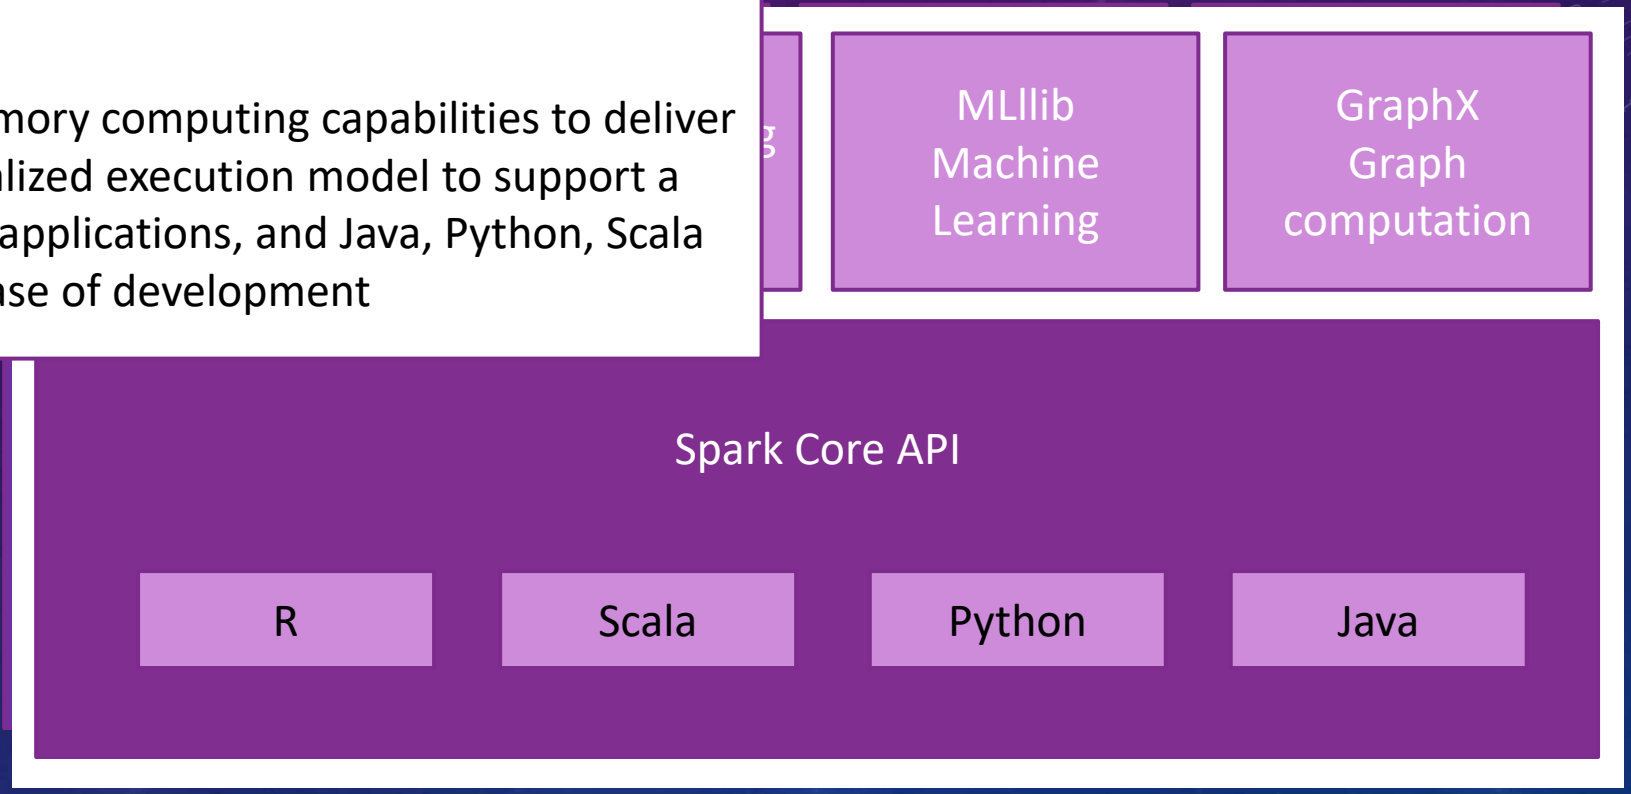

Spark Architecture

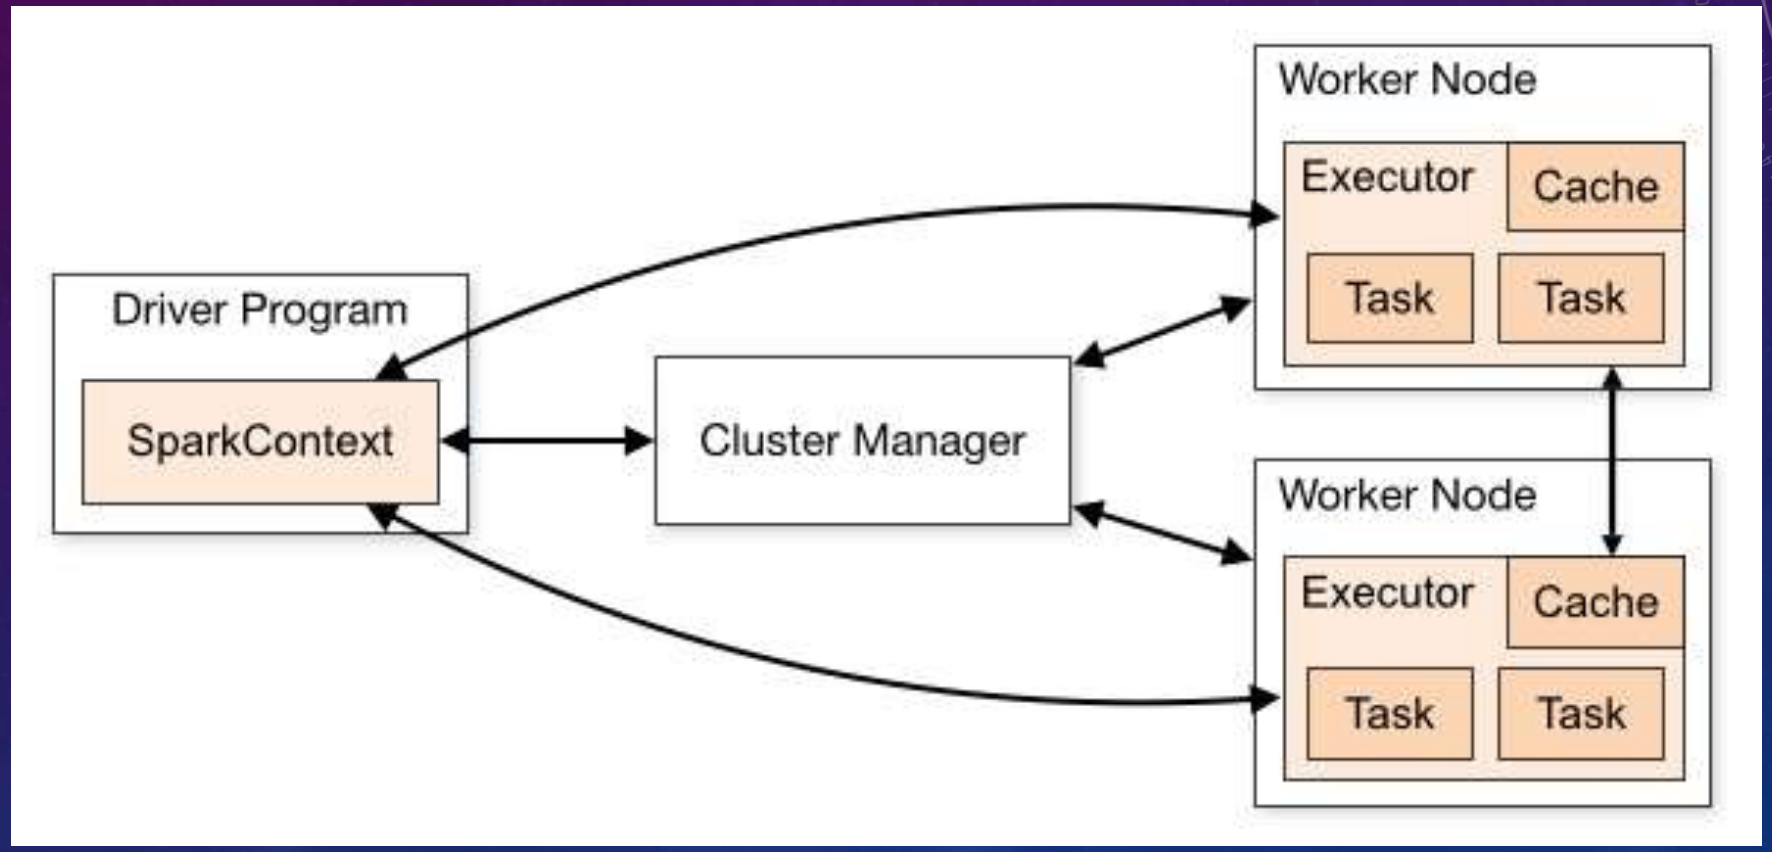

CONTEXT

+

Context

+

- Big Data analytics
- Faster

- Health care data increasing rapidly
- Problems in health care
  - High cost
  - High waste
  - Low quality

Context

Health care data volume is increasing and is expected to skyrocket as new ways of collecting health data in various forms are continuously emerging (i.e. patient centered e-health record, and wearables)

- Health care data increasing rapidly
- Problems in health care
  - High cost
  - High waste
  - Low quality

According to AbuKhousa and Campbell, the healthcare system has a massive wealth of information but knowledge poor, “ there is a lack of effectual analysis tools to discover knowledge contained in the databases of these systems ”

Context

+

Overall spending ($170 billion spend on health in 2015-16)

- Health care data increasing rapidly
- Problems in health care
  - High cost
  - High waste
  - Low quality

Context

+

Poor quality of health (increasing mortality and morbidity rate on preventable diseases)

- Health care data increasing rapidly
- Problems in health care
  - High cost
  - High waste
  - Low quality

PURPOSE

- Understand big data processing in health care
- Learn factors that contribute to heart disease
- Build machine learning model to predict heart disease.
- Use the Spark framework to implement analysis.
- Find significant risk factors of coronary heart disease

DELIVERABLES

- Build a predictive analytics platform to predict heart disease using Spark's Machine Learning library module in Standalone cluster mode with Mongo DB as the database.
- Determine which features or feature subsets contribute to risk of heart disease.
- Analyse result by comparing the ground truth to predicted outcome via graph or table.

OVERVIEW ON THE APPROACH USED

PROJECT MANAGEMENT APPROACH

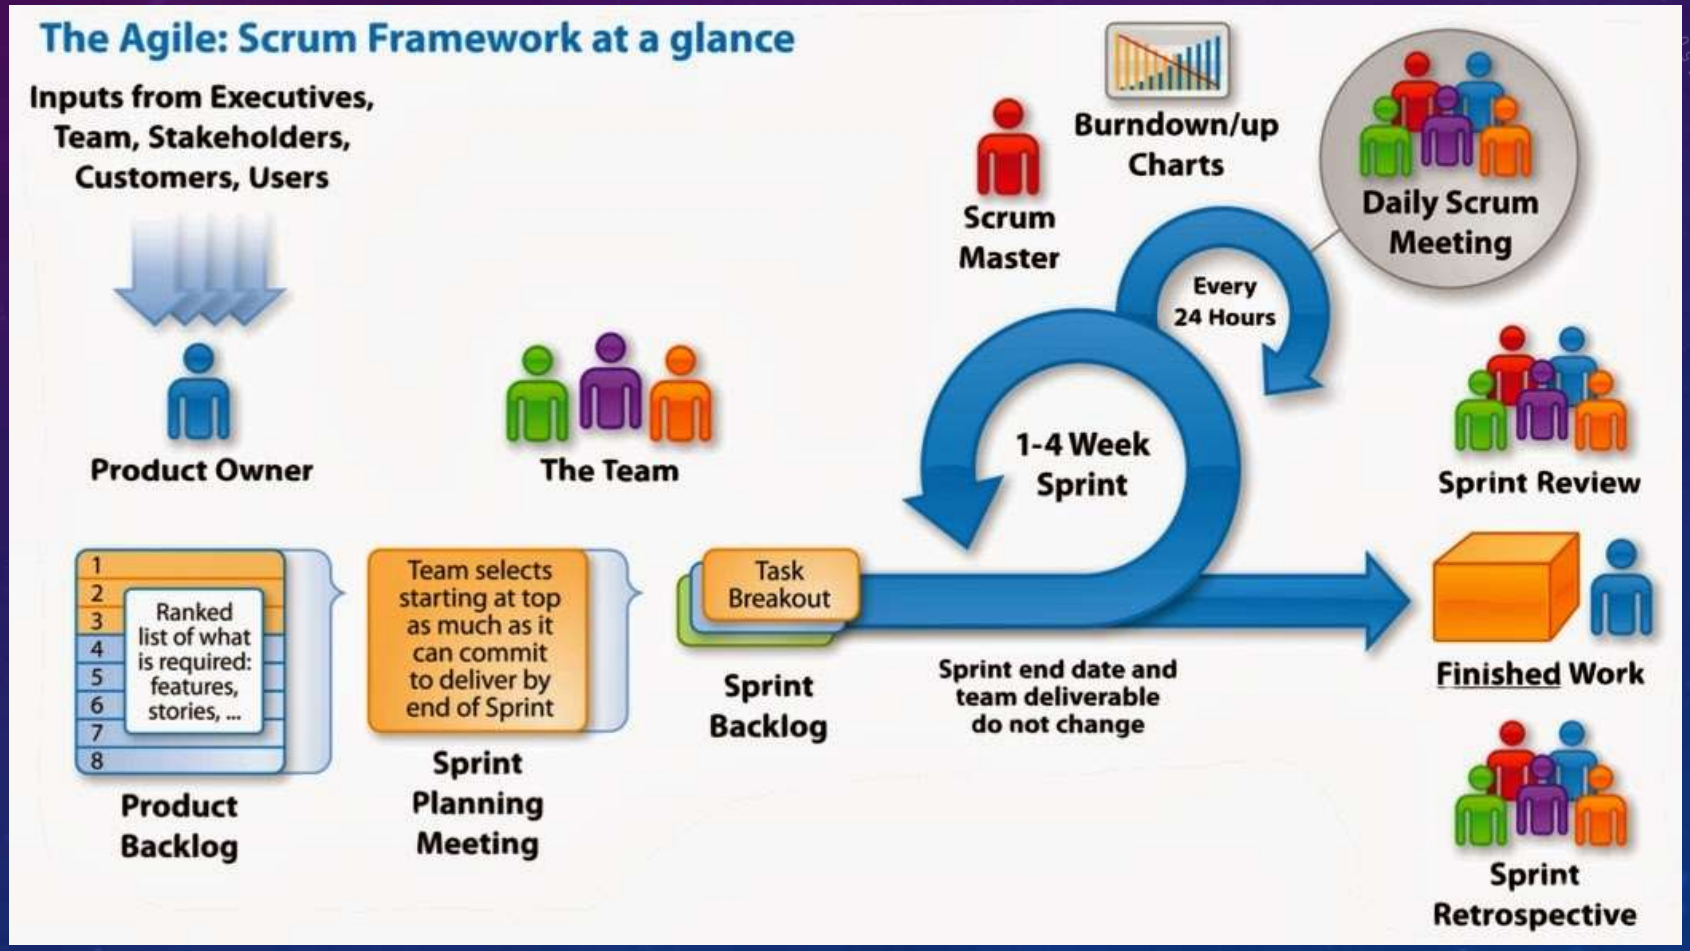

Project Architecture

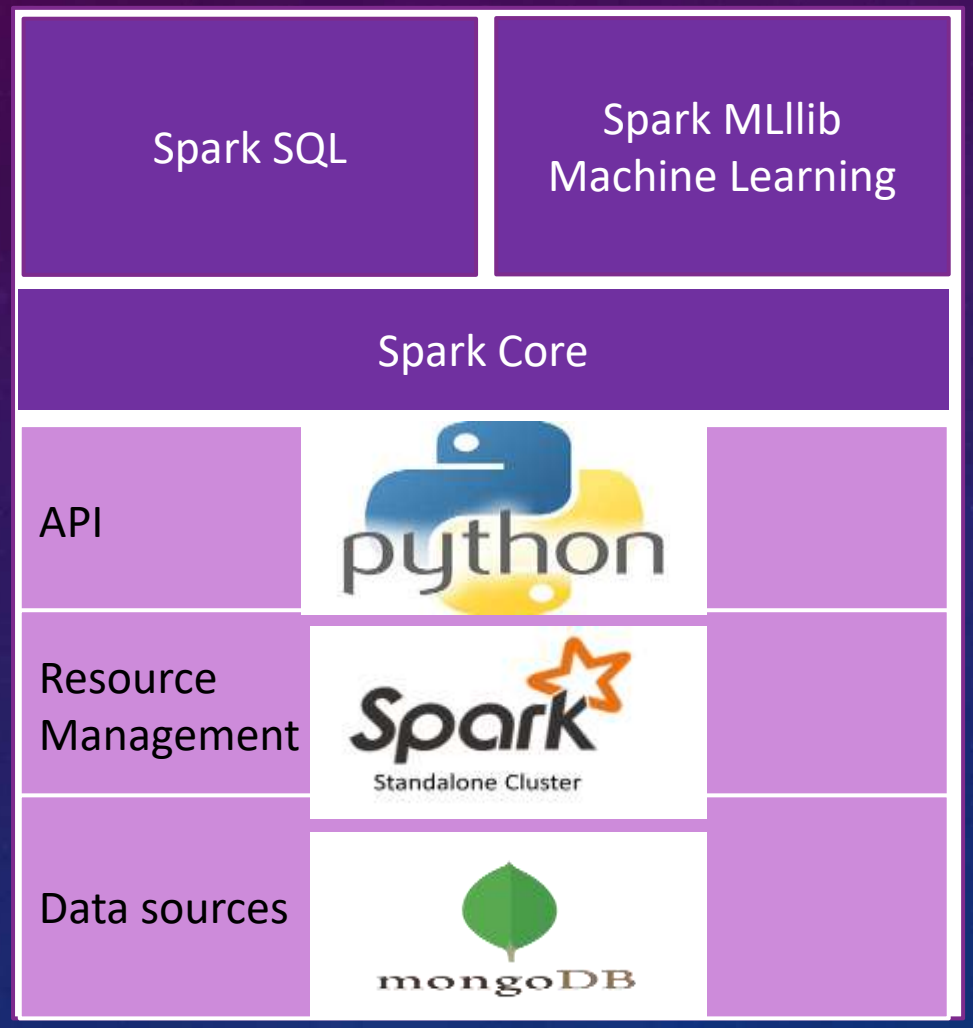

Project Approach/Methodology

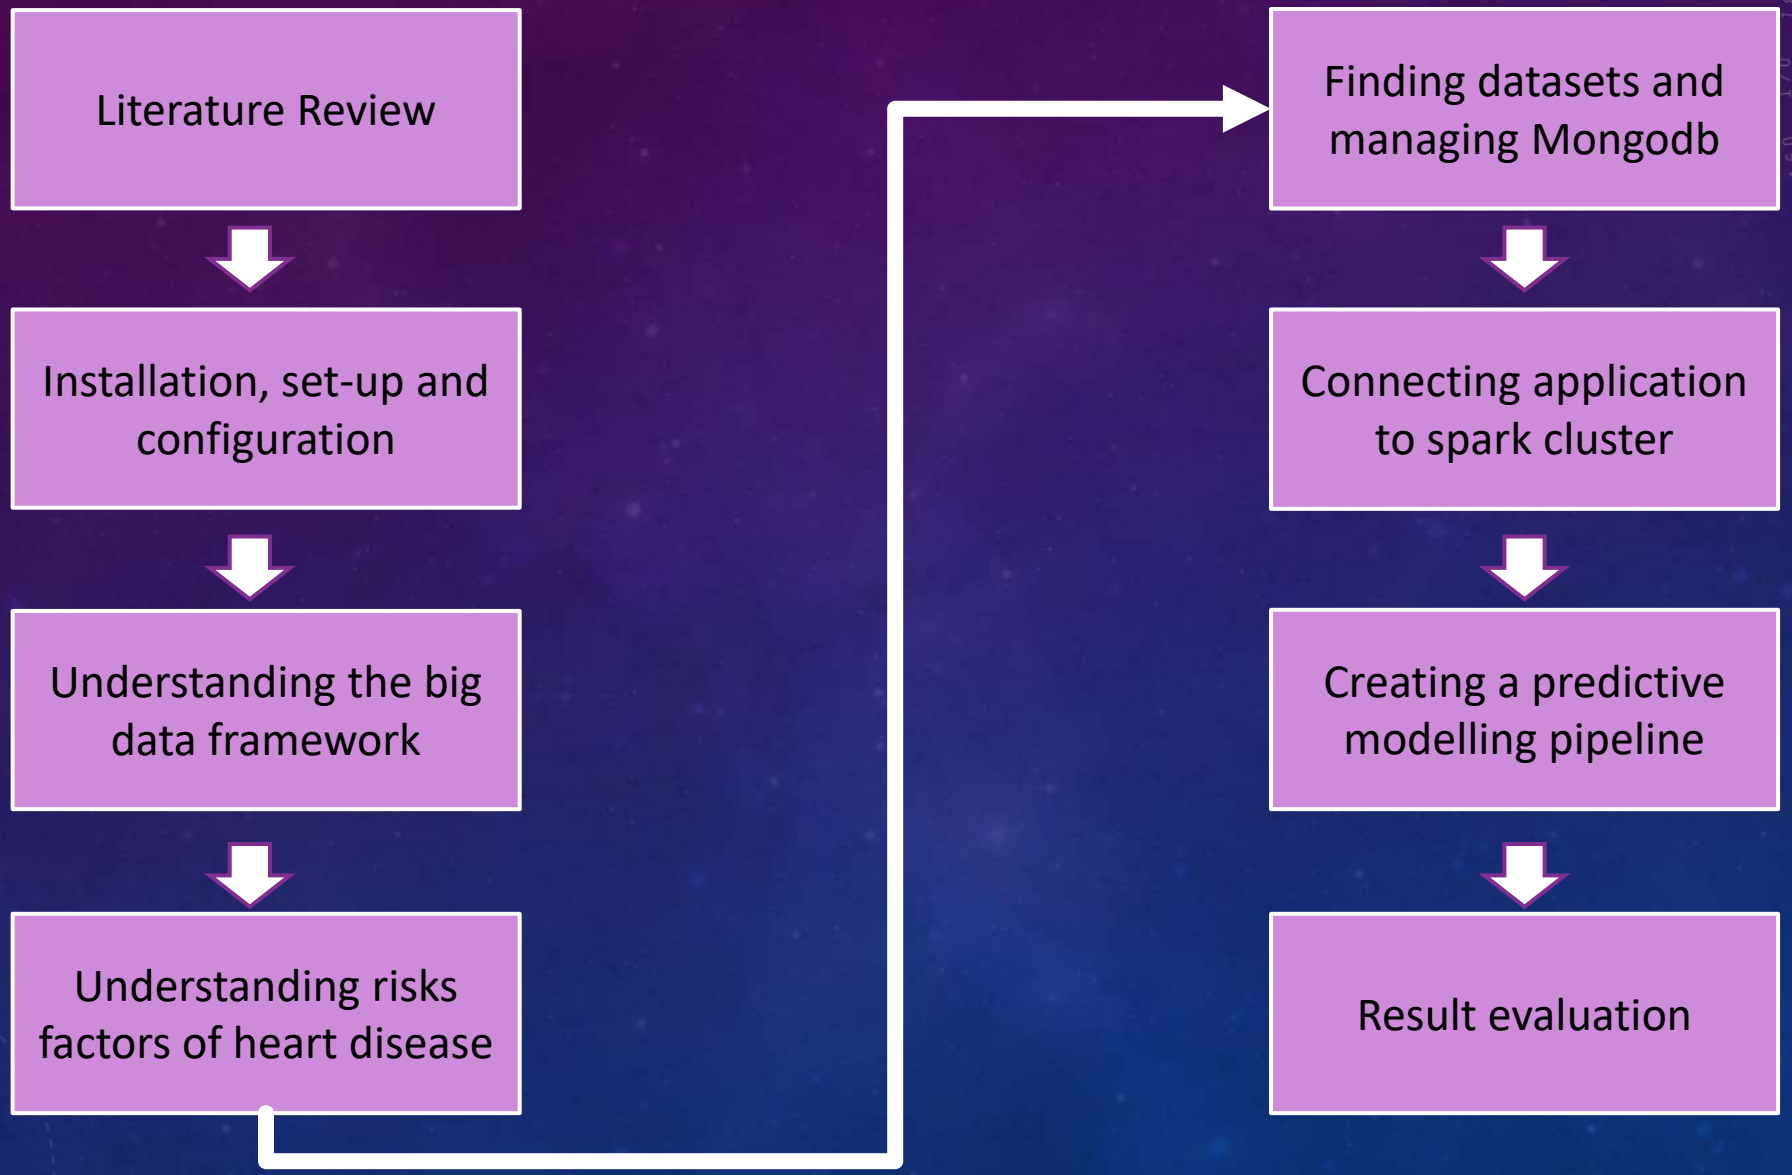

Project Approach/Methodology

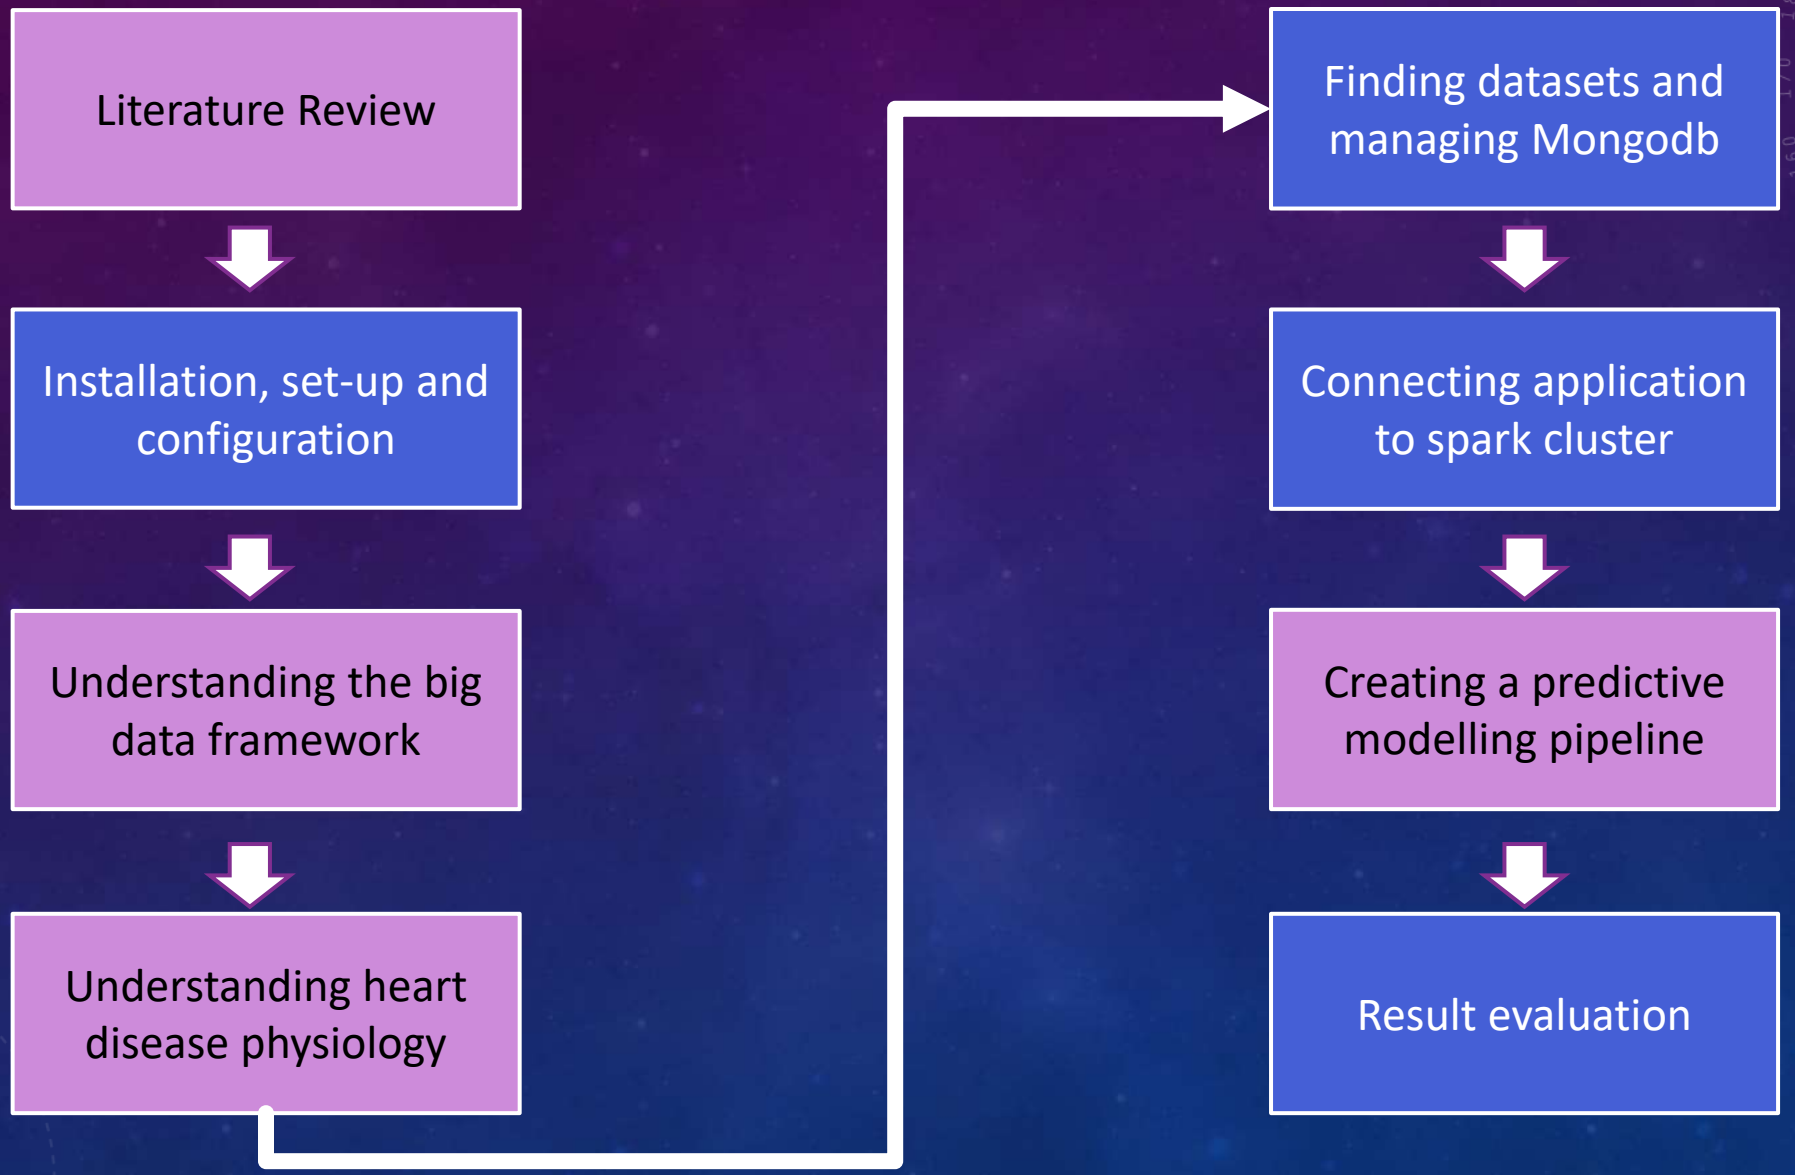

Project Approach/Methodology

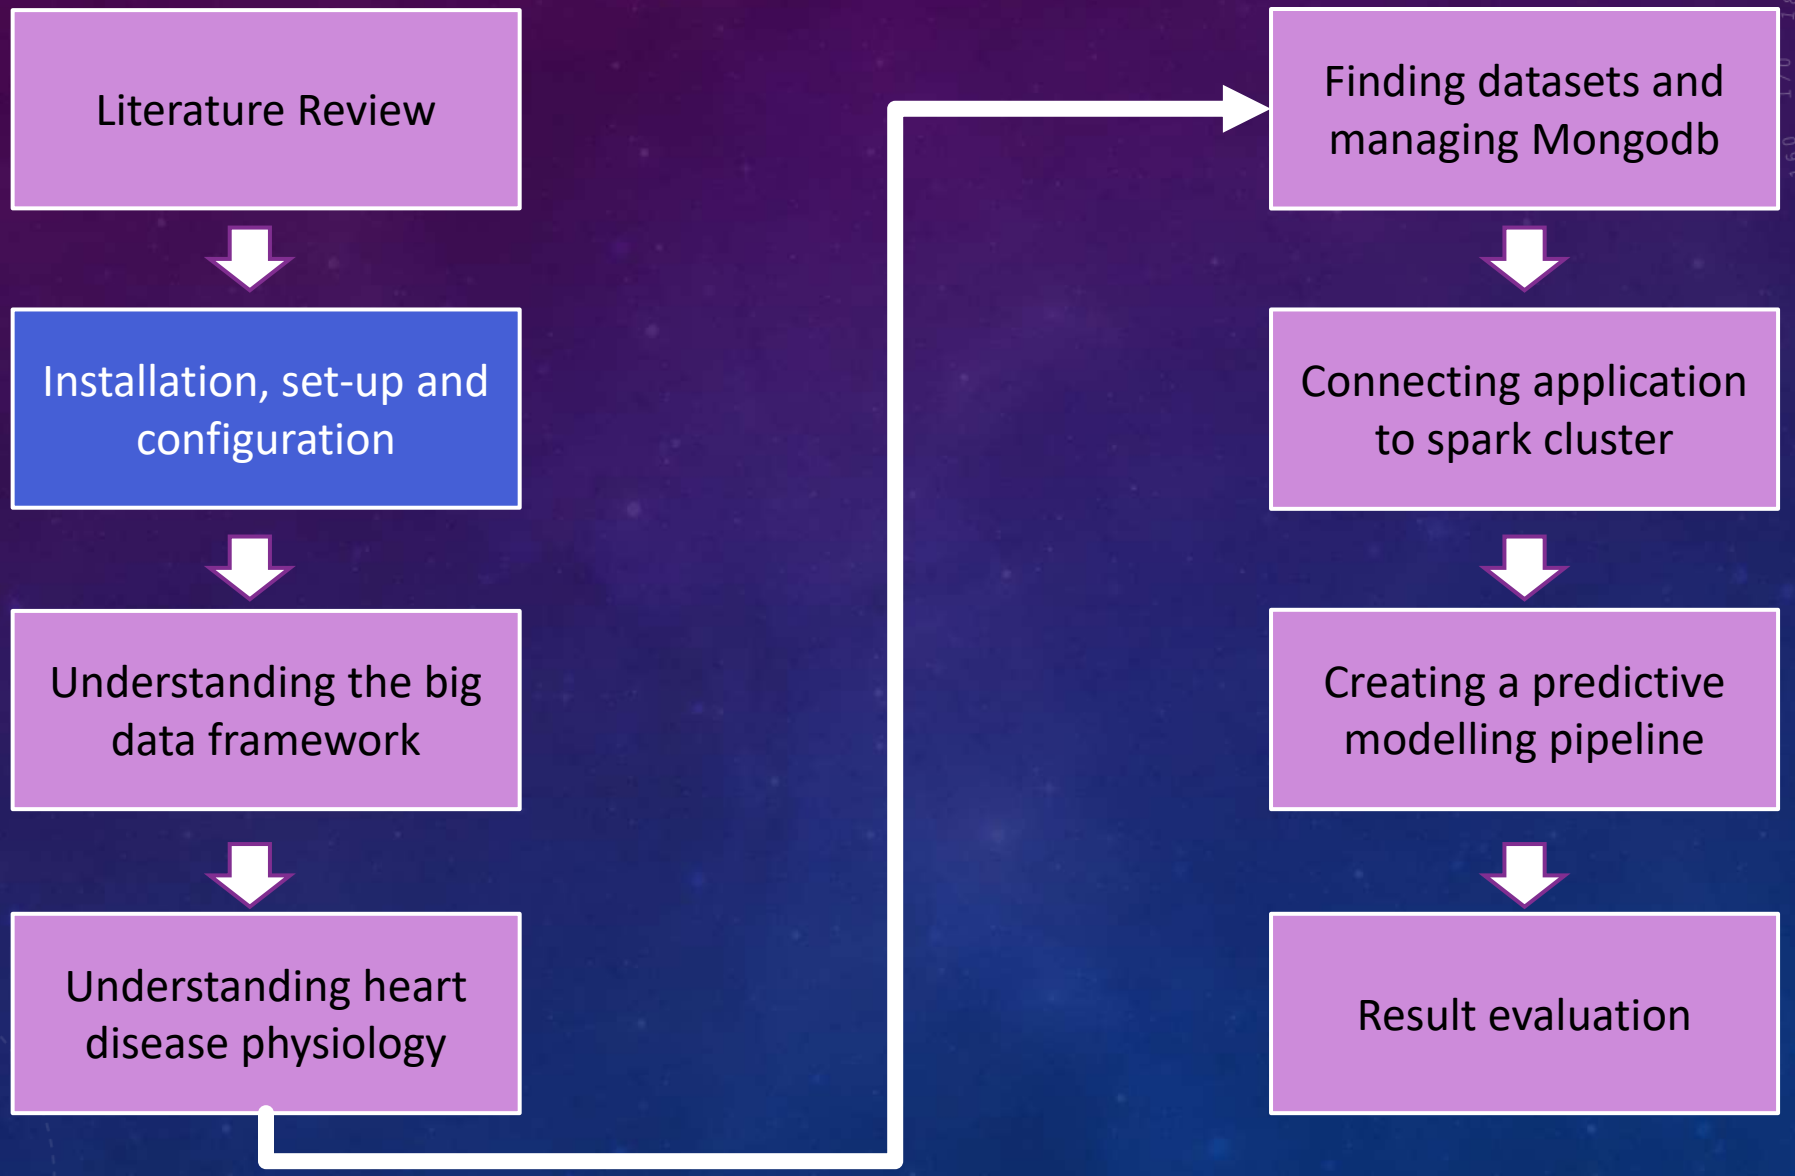

Installation, set-up and configuration of machine and tools

Virtual-Box/ Ubuntu setup

python 3.6.5

jupyter

Java Jdk8

Scala

- Install Py4j library
- Install pyspark

Download Spark 2.3.1

Setup environmental variables/path

mongoDB

- Install pymongo
- Install mongo-spark connector

Project Approach/Methodology

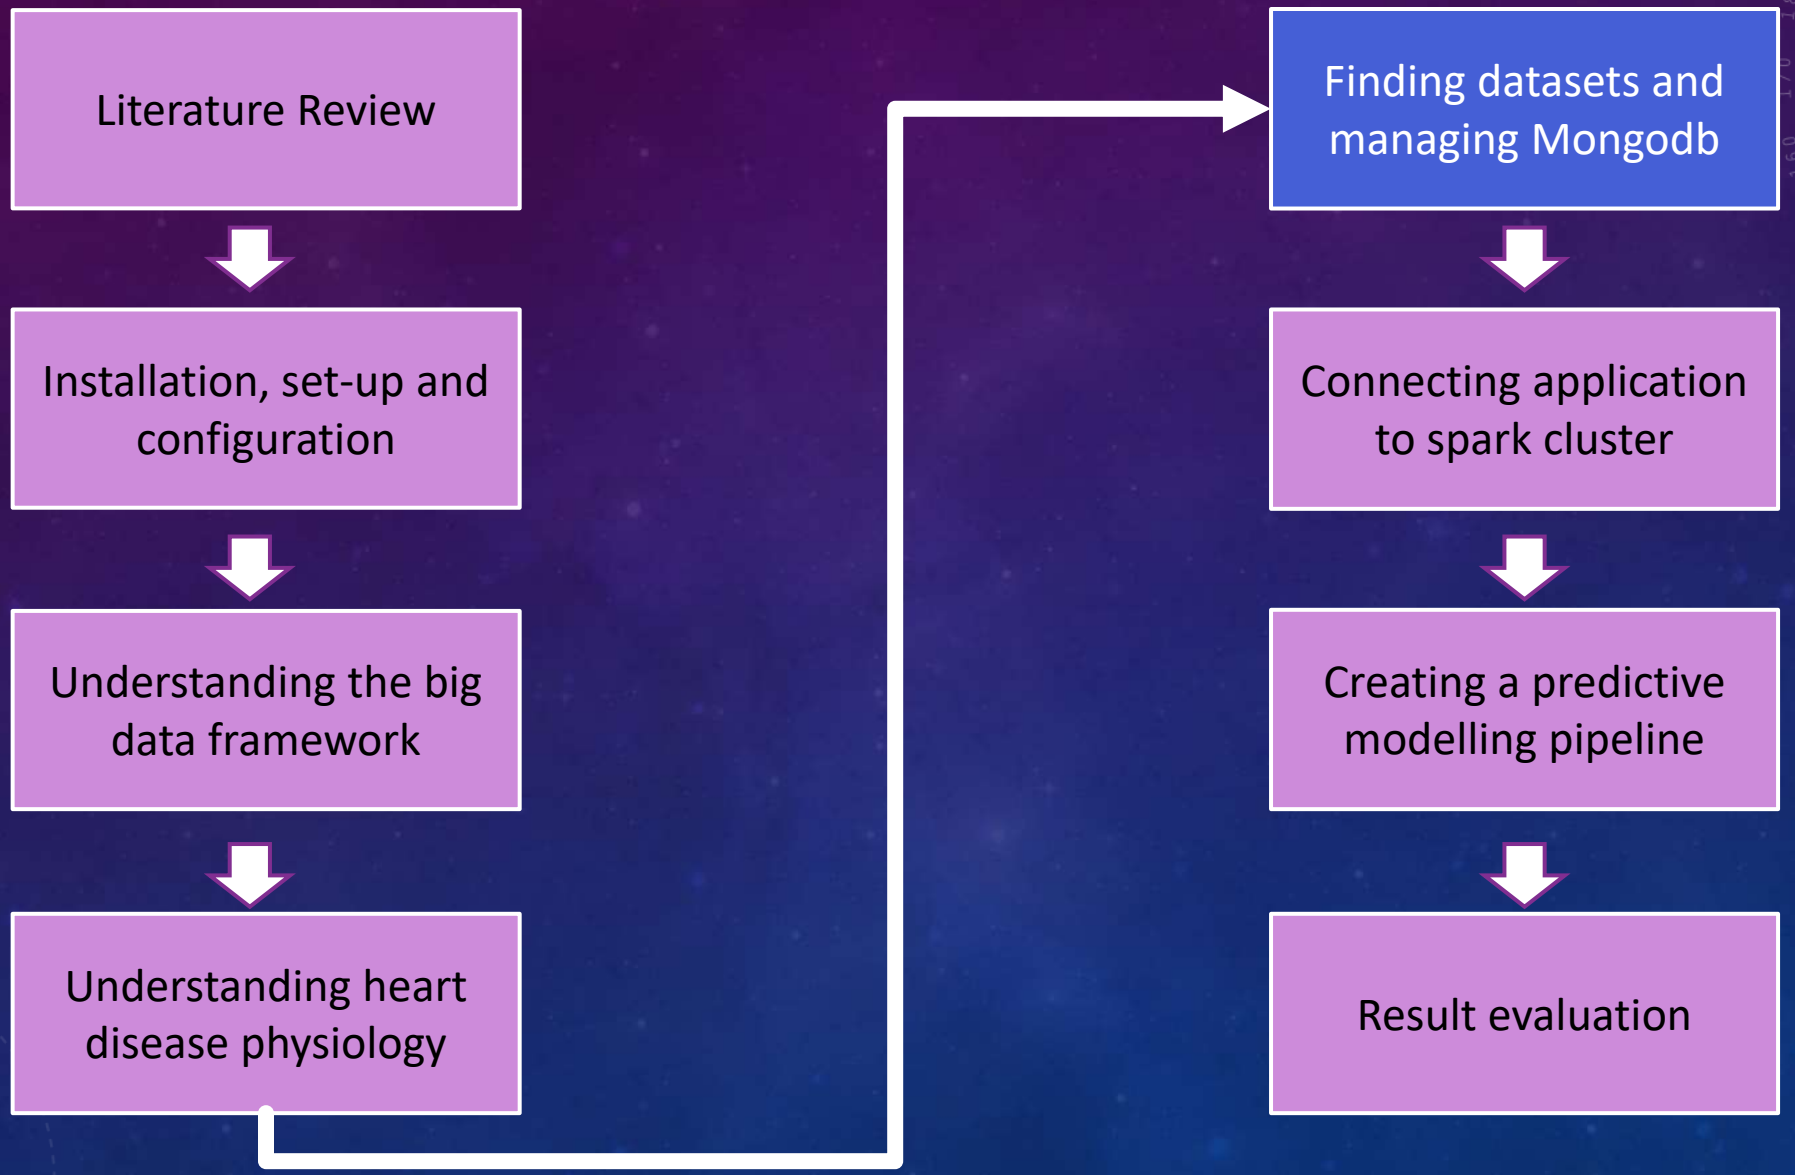

Project Approach/Methodology: Datasets and MongoDB

Framingham Heart Study datasets

	age	sex	education	currentSmoker	cigsPerDay	heartRate	BMI	glucose	diabetes	sysBP	diaBP	BPMeds	prevalentHyp	prevalentStroke	totChol	T
0	40	Male	3	Yes	70	98	31.57	80	No	132.0	86.0	No	Yes	No	210.0	
1	56	Male	1	Yes	60	70	29.64	85	No	125.0	79.0	No	No	No	246.0	
2	59	Male	1	Yes	60	70	25.05	84	No	153.5	105.0	No	Yes	No	298.0	
3	58	Male	2	Yes	60	75	32.00	65	No	150.0	97.0	No	Yes	No	250.0	
4	39	Male	1	Yes	60	59	23.60	78	No	112.0	65.0	No	No	No	215.0	

- csv format
- 4241 records

Project Approach/Methodology: Datasets and MongoDB

Framingham Heart Study datasets

Categorical variables	Numeric/continuous variables
Sex	age
Education	cigsPerDay
currentSmoker	heartrate
Diabetes	BMI
BPMeds	Glucose
prevalentHyp	sysBP
prevalentStroke	diaBP
label	totChol

Project Approach/Methodology: Managing MongoDB

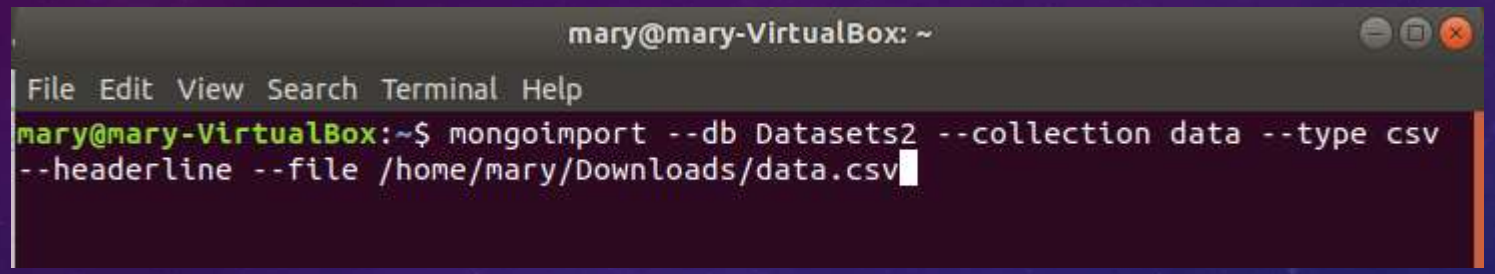

- Importing datasets

Project Approach/Methodology: Managing MongoDB

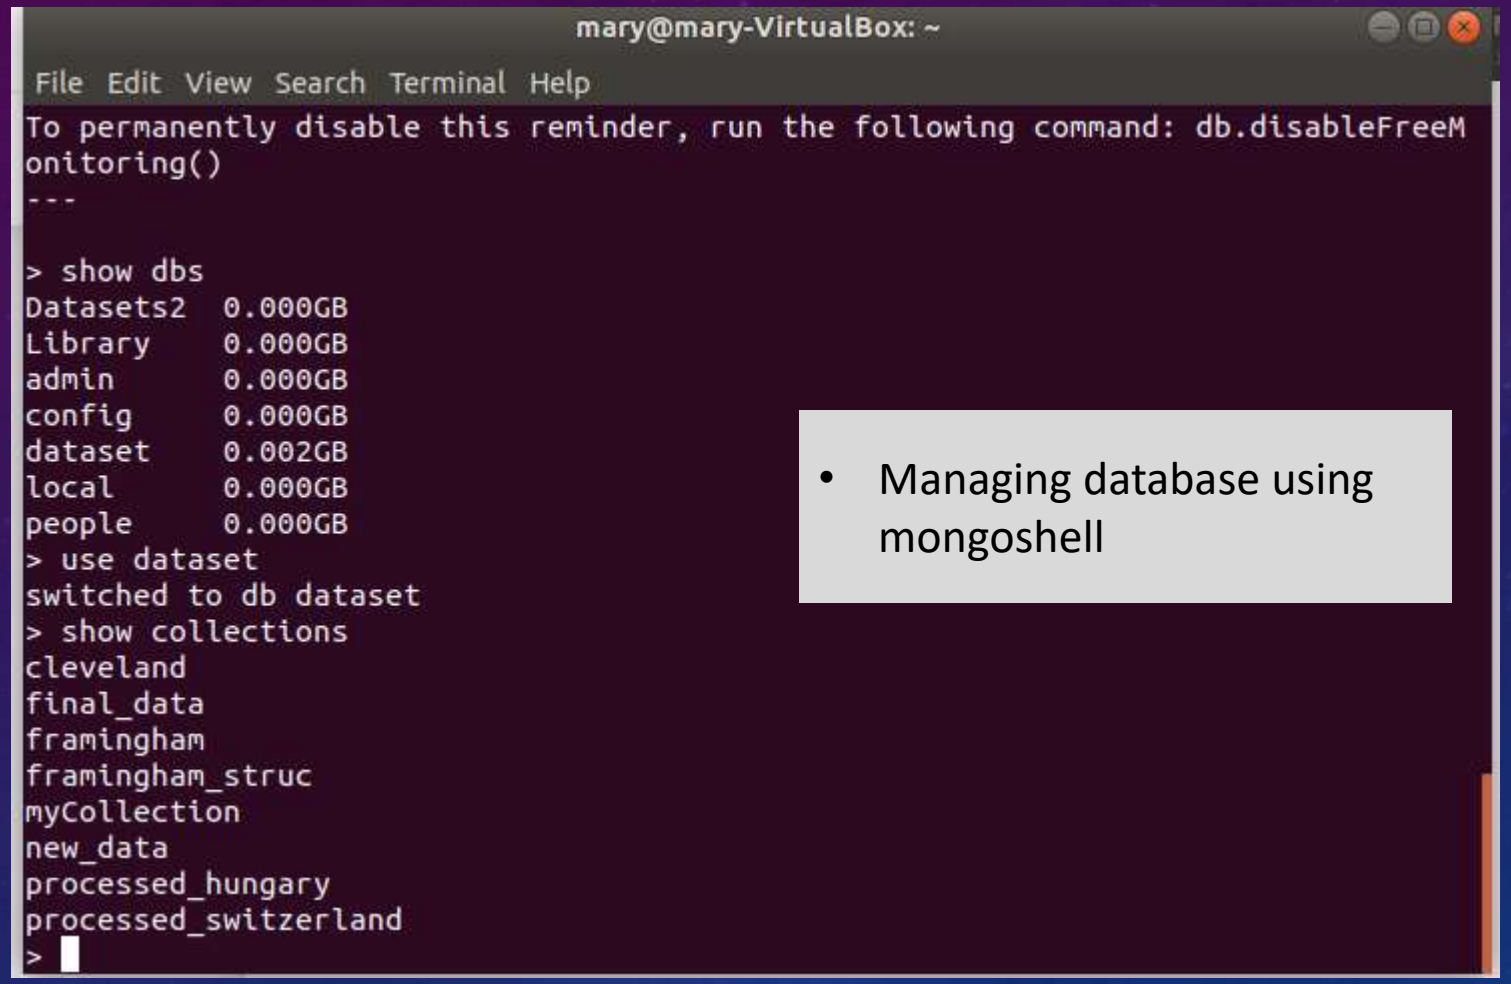

- Managing database using mongoshell

Project Approach/Methodology

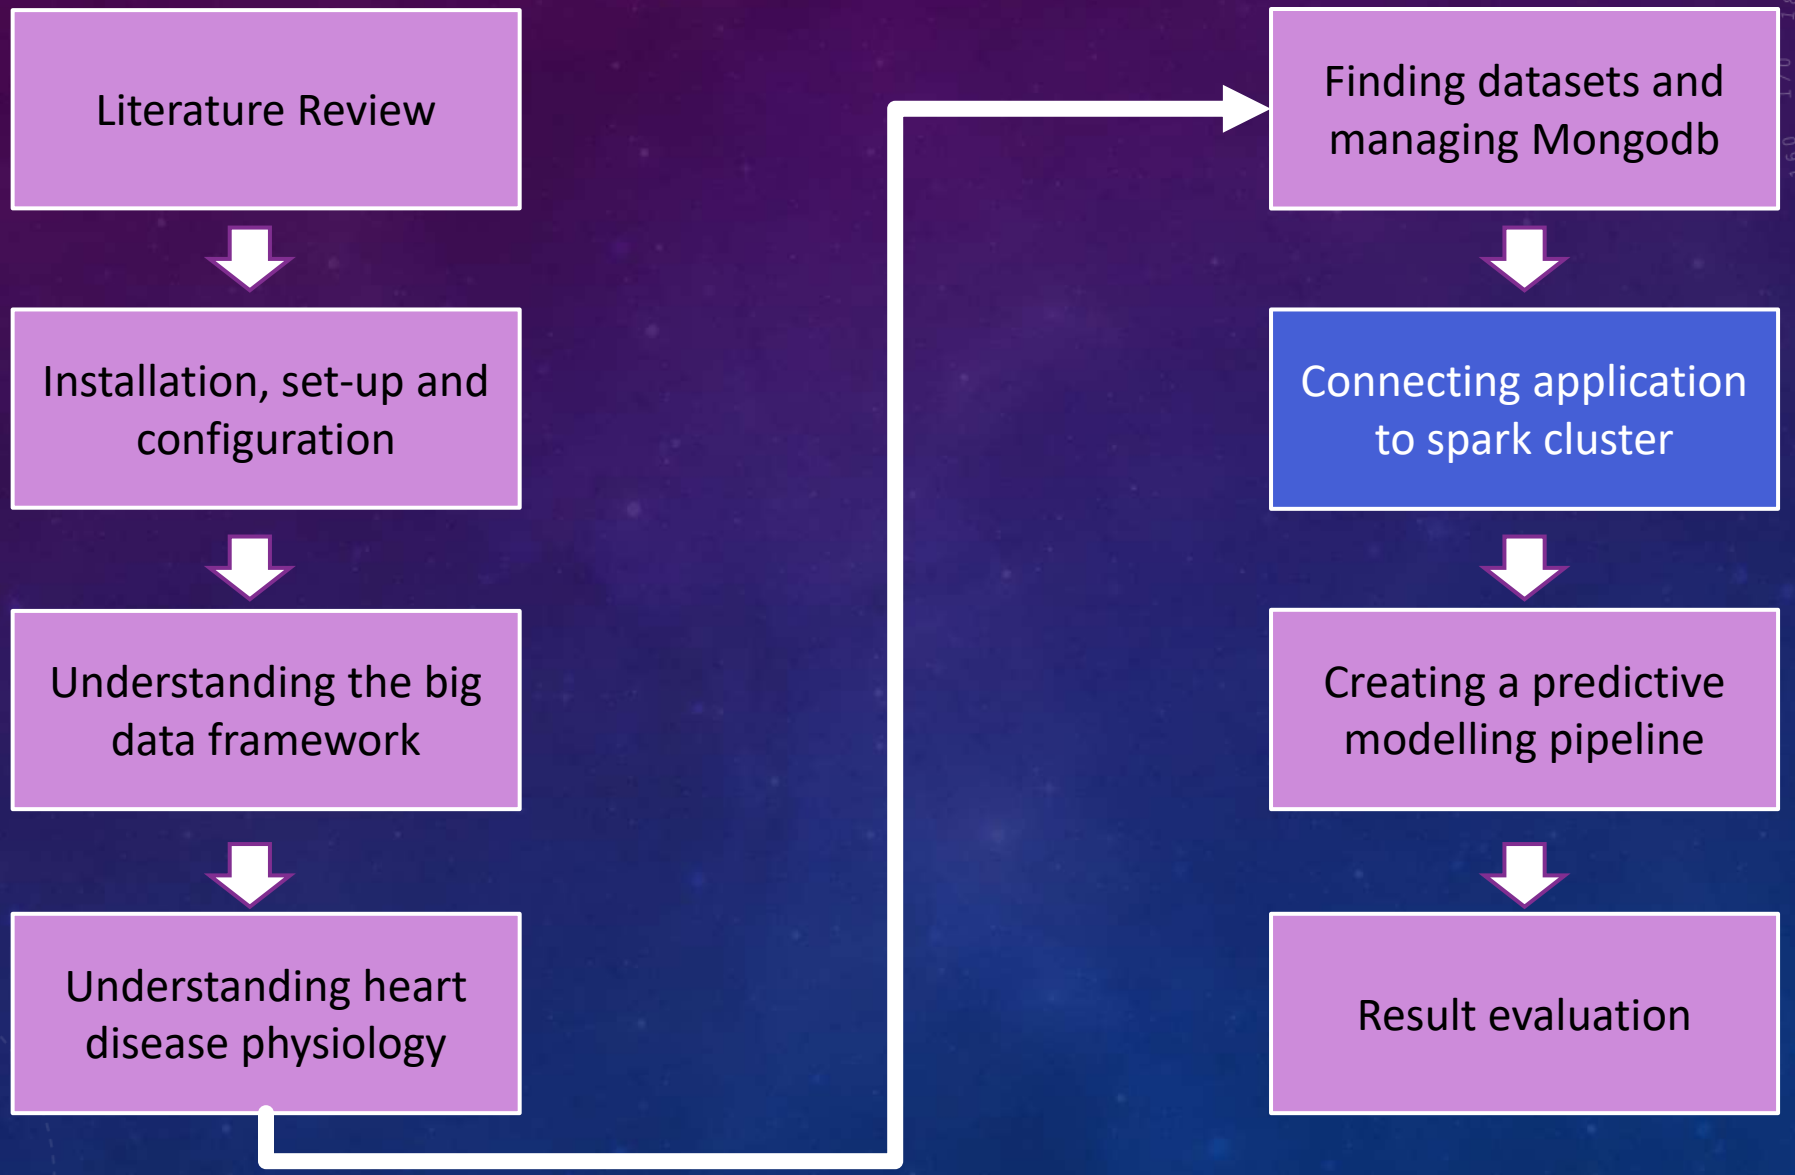

Deployment stage: Starting Mongodb service

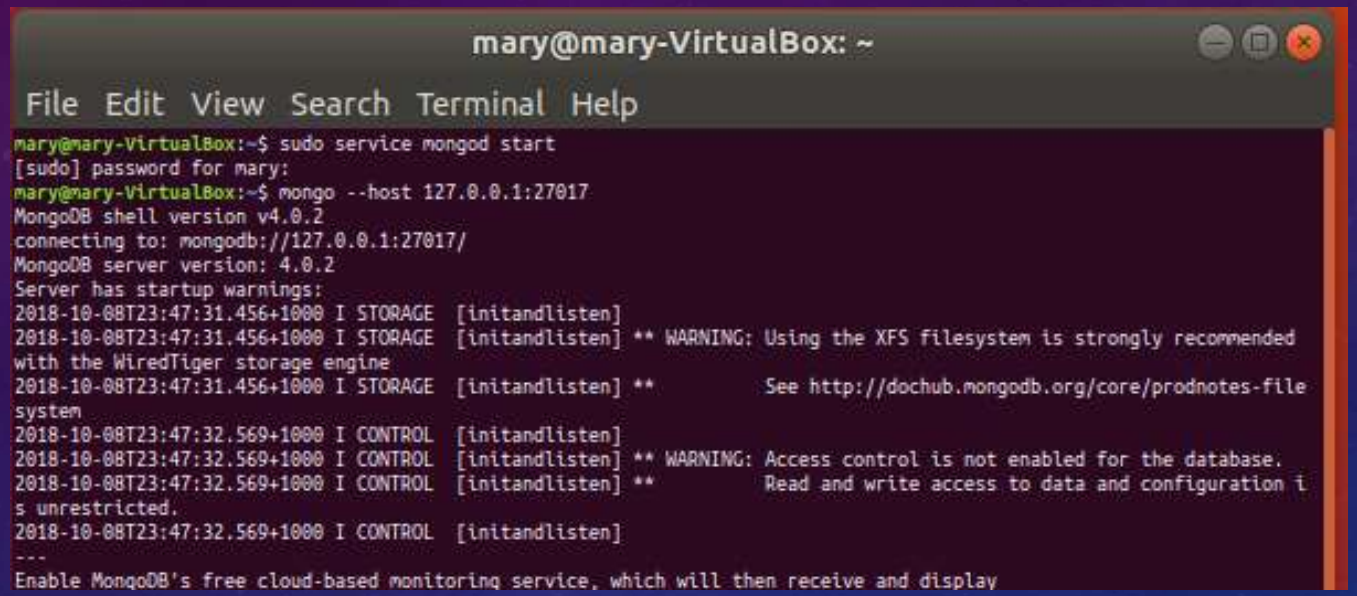

Starting Mongodb service

Deployment stage : starting the spark master cluster

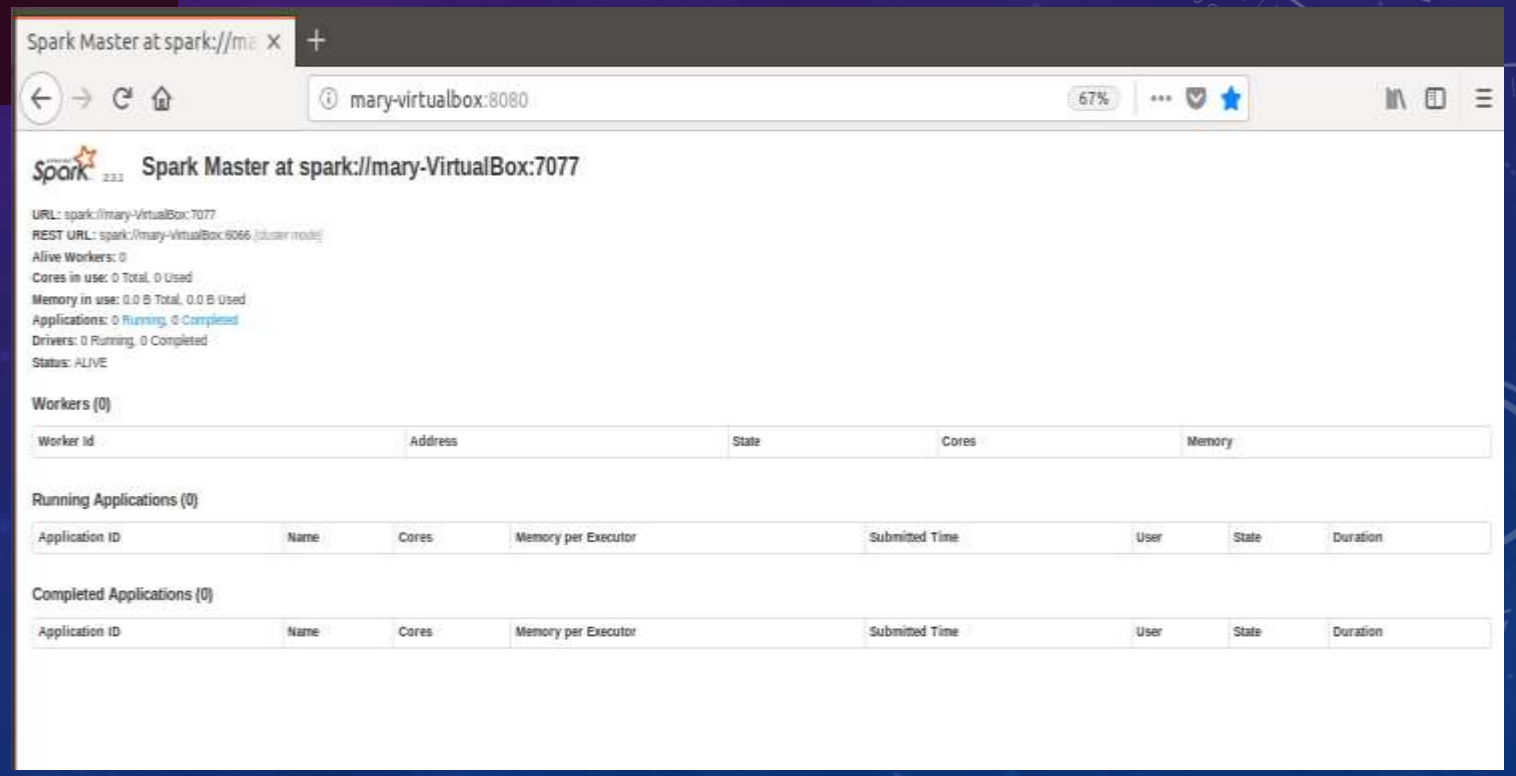

Deployment stage: starting 4 worker nodes

```
mary@mary-VirtualBox: ~/spark-2.3.1
File  Edit  View  Search  Terminal  Help
mary@mary-VirtualBox:~/spark-2.3.1$ ./sbin/start-slave.sh spark://mary-VirtualBo
x:7077
starting org.apache.spark.deploy.worker.Worker, logging to /home/mary/spark-2.3.
1/logs/spark-mary-org.apache.spark.deploy.worker.Worker-1-mary-VirtualBox.out
starting org.apache.spark.deploy.worker.Worker, logging to /home/mary/spark-2.3.
1/logs/spark-mary-org.apache.spark.deploy.worker.Worker-2-mary-VirtualBox.out
starting org.apache.spark.deploy.worker.Worker, logging to /home/mary/spark-2.3.
1/logs/spark-mary-org.apache.spark.deploy.worker.Worker-3-mary-VirtualBox.out
starting org.apache.spark.deploy.worker.Worker, logging to /home/mary/spark-2.3.
1/logs/spark-mary-org.apache.spark.deploy.worker.Worker-4-mary-VirtualBox.out
mary@mary-VirtualBox:~/spark-2.3.1$
mary@mary-VirtualBox:~/spark-2.3.1$
mary@mary-VirtualBox:~/spark-2.3.1$
mary@mary-VirtualBox:~/spark-2.3.1$
mary@mary-VirtualBox:~/spark-2.3.1$
mary@mary-VirtualBox:~/spark-2.3.1$
mary@mary-VirtualBox:~/spark-2.3.1$
mary@mary-VirtualBox:~/spark-2.3.1$
```

Starting the Spark cluster

Spark Master at spark://ma

mary-virtualbox:8080

67%

Spark Master at spark://mary-VirtualBox:7077

URL: spark://mary-VirtualBox:7077
REST URL: spark://mary-VirtualBox:6066 *(cluster mode)*
Alive Workers: 4
Cores in use: 4 Total, 0 Used
Memory in use: 8.0 GB Total, 0.0 B Used
Applications: 0 Running, 0 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers (4)

Worker Id	Address	State	Cores	Memory
worker-20181008235626-10.0.2.15-46867	10.0.2.15:46867	ALIVE	1 (0 Used)	2.0 GB (0.0 B Used)
worker-20181008235629-10.0.2.15-37025	10.0.2.15:37025	ALIVE	1 (0 Used)	2.0 GB (0.0 B Used)
worker-20181008235631-10.0.2.15-39867	10.0.2.15:39867	ALIVE	1 (0 Used)	2.0 GB (0.0 B Used)
worker-20181008235634-10.0.2.15-38437	10.0.2.15:38437	ALIVE	1 (0 Used)	2.0 GB (0.0 B Used)

Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

Deployment stage: starting 4 worker nodes

jupyter SparkDemo_DecisionTrees-specific Last Checkpoint: Last Friday at 12:52 PM (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help

Trusted | Python 3

Run Code

Decision Tree Methods Framingham dataset

Start Spark Session and connect application to cluster

In [*]:

```
from pyspark.sql import SparkSession
spark = SparkSession.builder \
    .master("spark://mary-VirtualBox:7077") \
    .appName('SparkDemo_DecisionTrees_FraminghamDataset_specific') \
    .config("spark.mongodb.input.uri=mongodb://127.0.0.1:27017/localhost:27017/dataset.framingham?readPref
    .config("spark.mongodb.output.uri=mongodb://127.0.0.1:27017/localhost:27017/dataset.framingham") \
    .getOrCreate()
```

Application running

SparkDemo_DecisionTre

67%

Spark 2.3.1 **Spark Master at spark://mary-VirtualBox:7077**

URL: spark://mary-VirtualBox:7077
REST URL: spark://mary-VirtualBox:6066 *(cluster mode)*
Alive Workers: 4
Cores in use: 4 Total, 4 Used
Memory in use: 8.0 GB Total, 4.0 GB Used
Applications: 1 Running, 0 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers (4)

Worker Id	Address	State	Cores	Memory
worker-20181008235626-10.0.2.15-46867	10.0.2.15:46867	ALIVE	1 (1 Used)	2.0 GB (1024.0 MB Used)
worker-20181008235629-10.0.2.15-37025	10.0.2.15:37025	ALIVE	1 (1 Used)	2.0 GB (1024.0 MB Used)
worker-20181008235631-10.0.2.15-39867	10.0.2.15:39867	ALIVE	1 (1 Used)	2.0 GB (1024.0 MB Used)
worker-20181008235634-10.0.2.15-38437	10.0.2.15:38437	ALIVE	1 (1 Used)	2.0 GB (1024.0 MB Used)

Running Applications (1)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
app-20181009000358-0000 (kill)	SparkDemo_DecisionTrees_FraminghamDataset_specific	4	1024.0 MB	2018/10/09 00:03:58	mary	RUNNING	50 s

Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

Deployment stage:

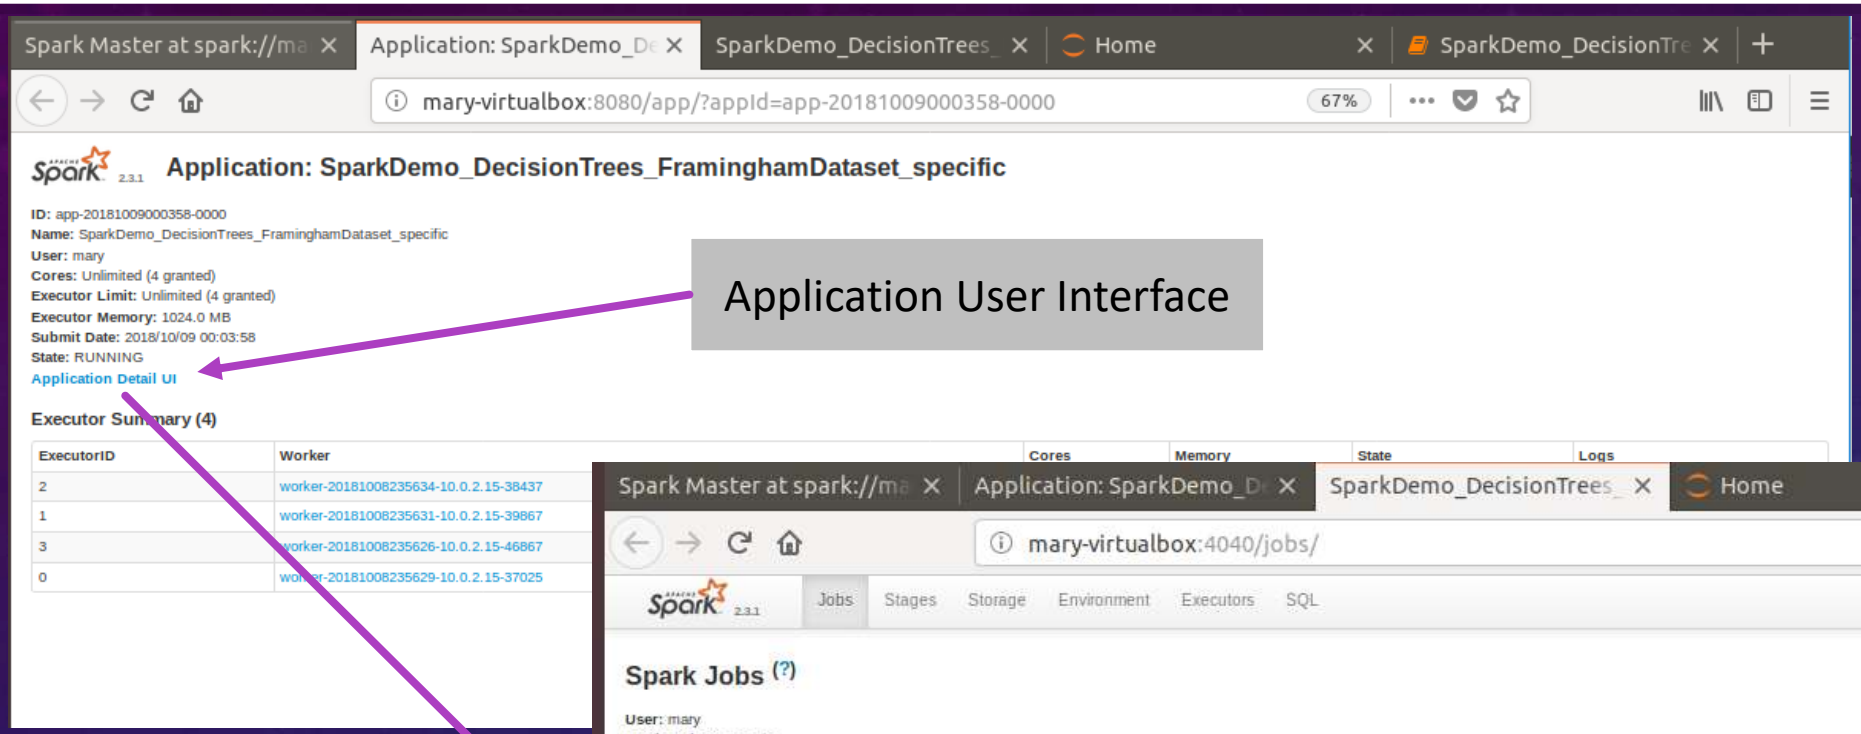

Application User Interface

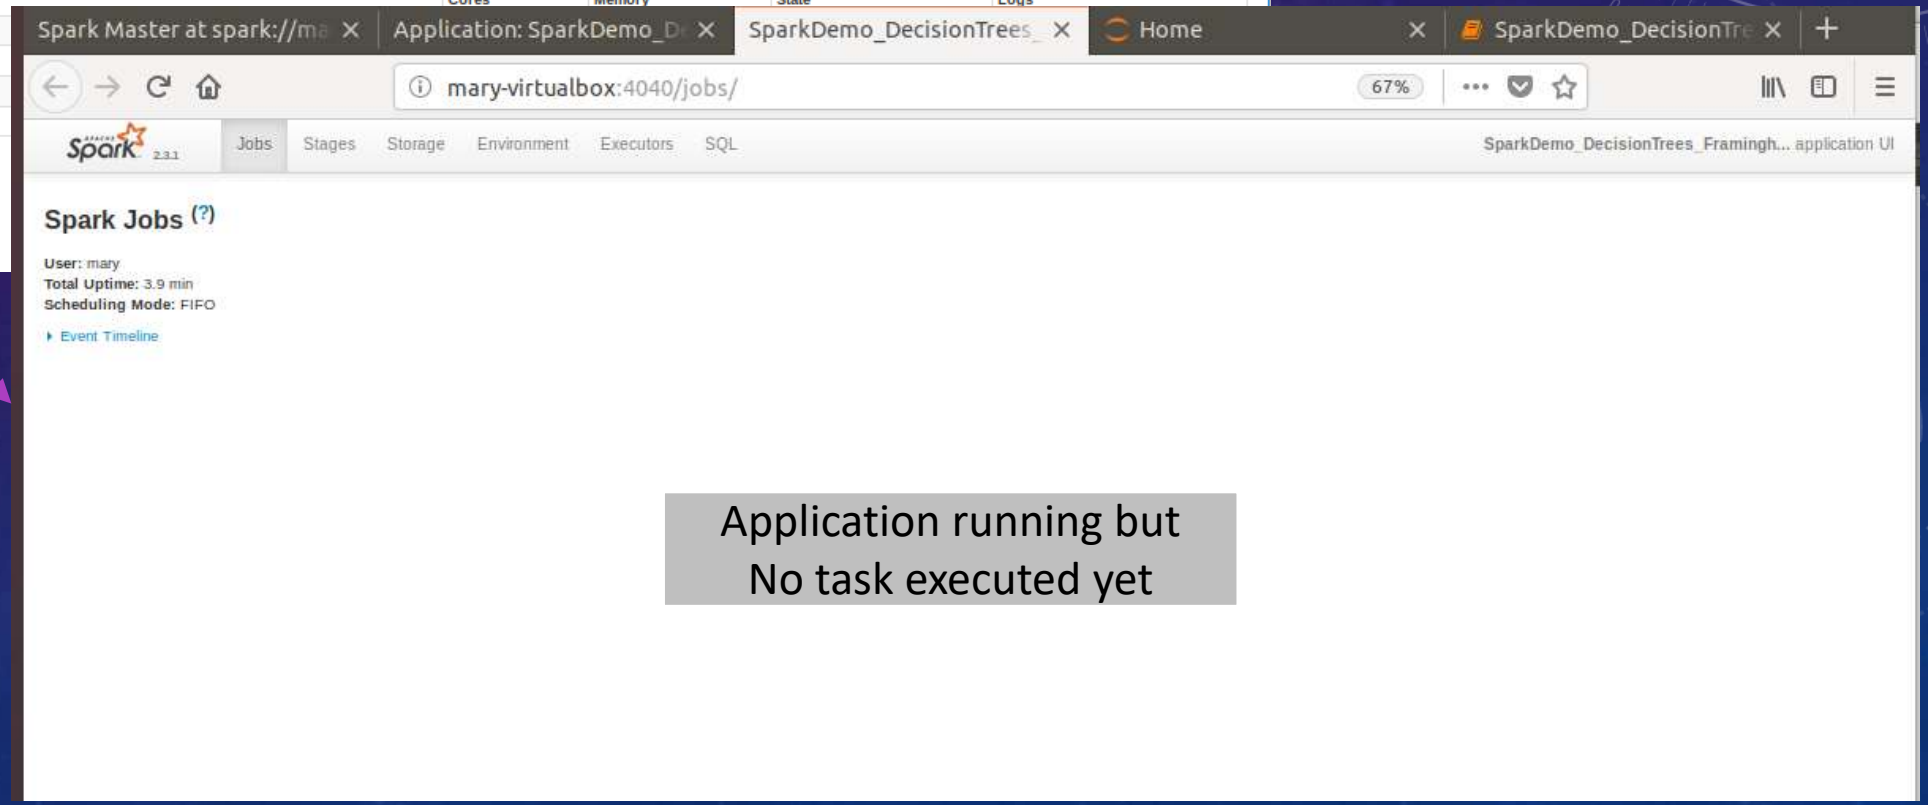

Application running but
No task executed yet

Deployment stage:

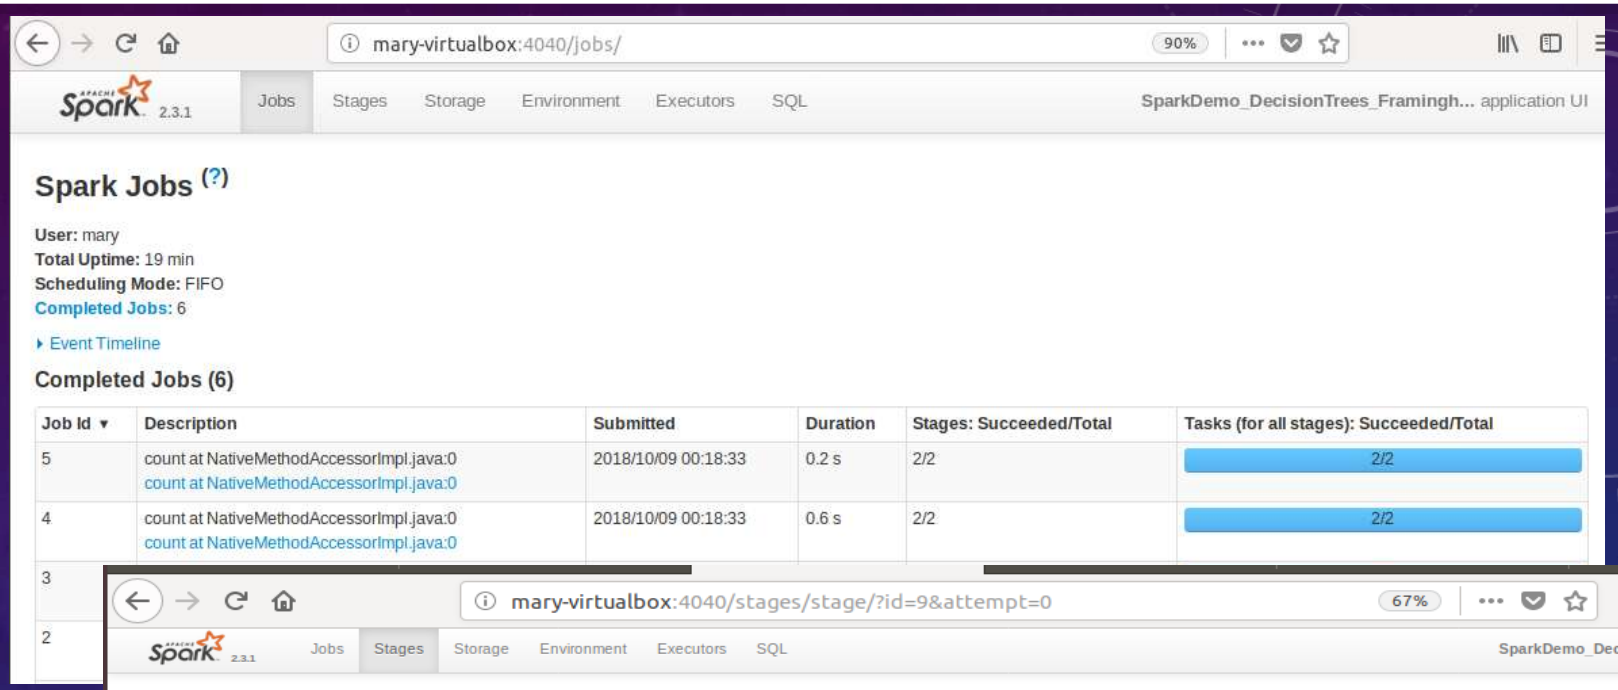

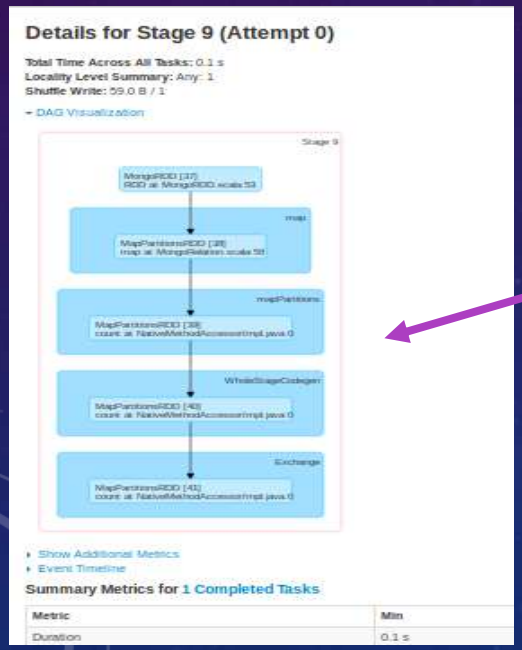

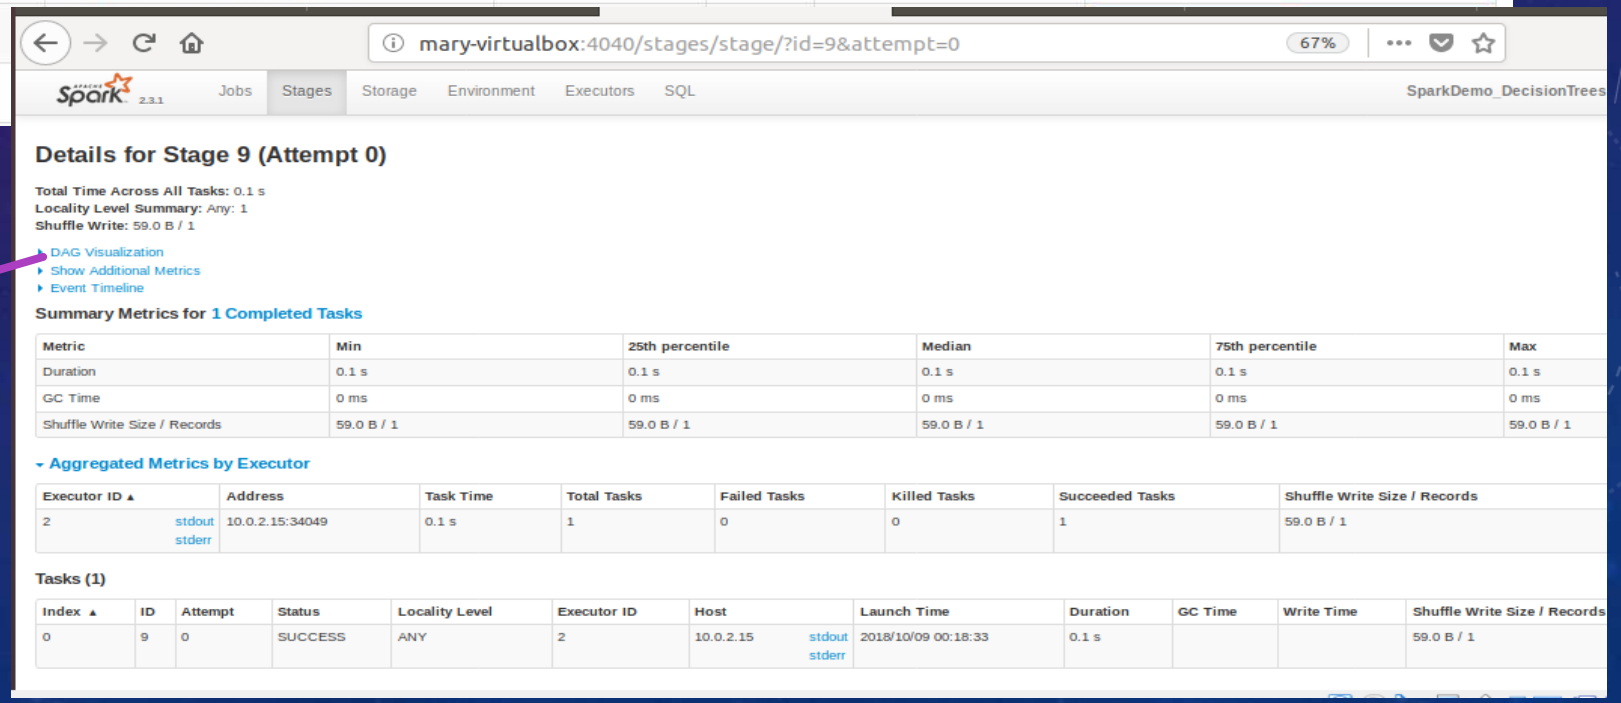

Deployment stage:

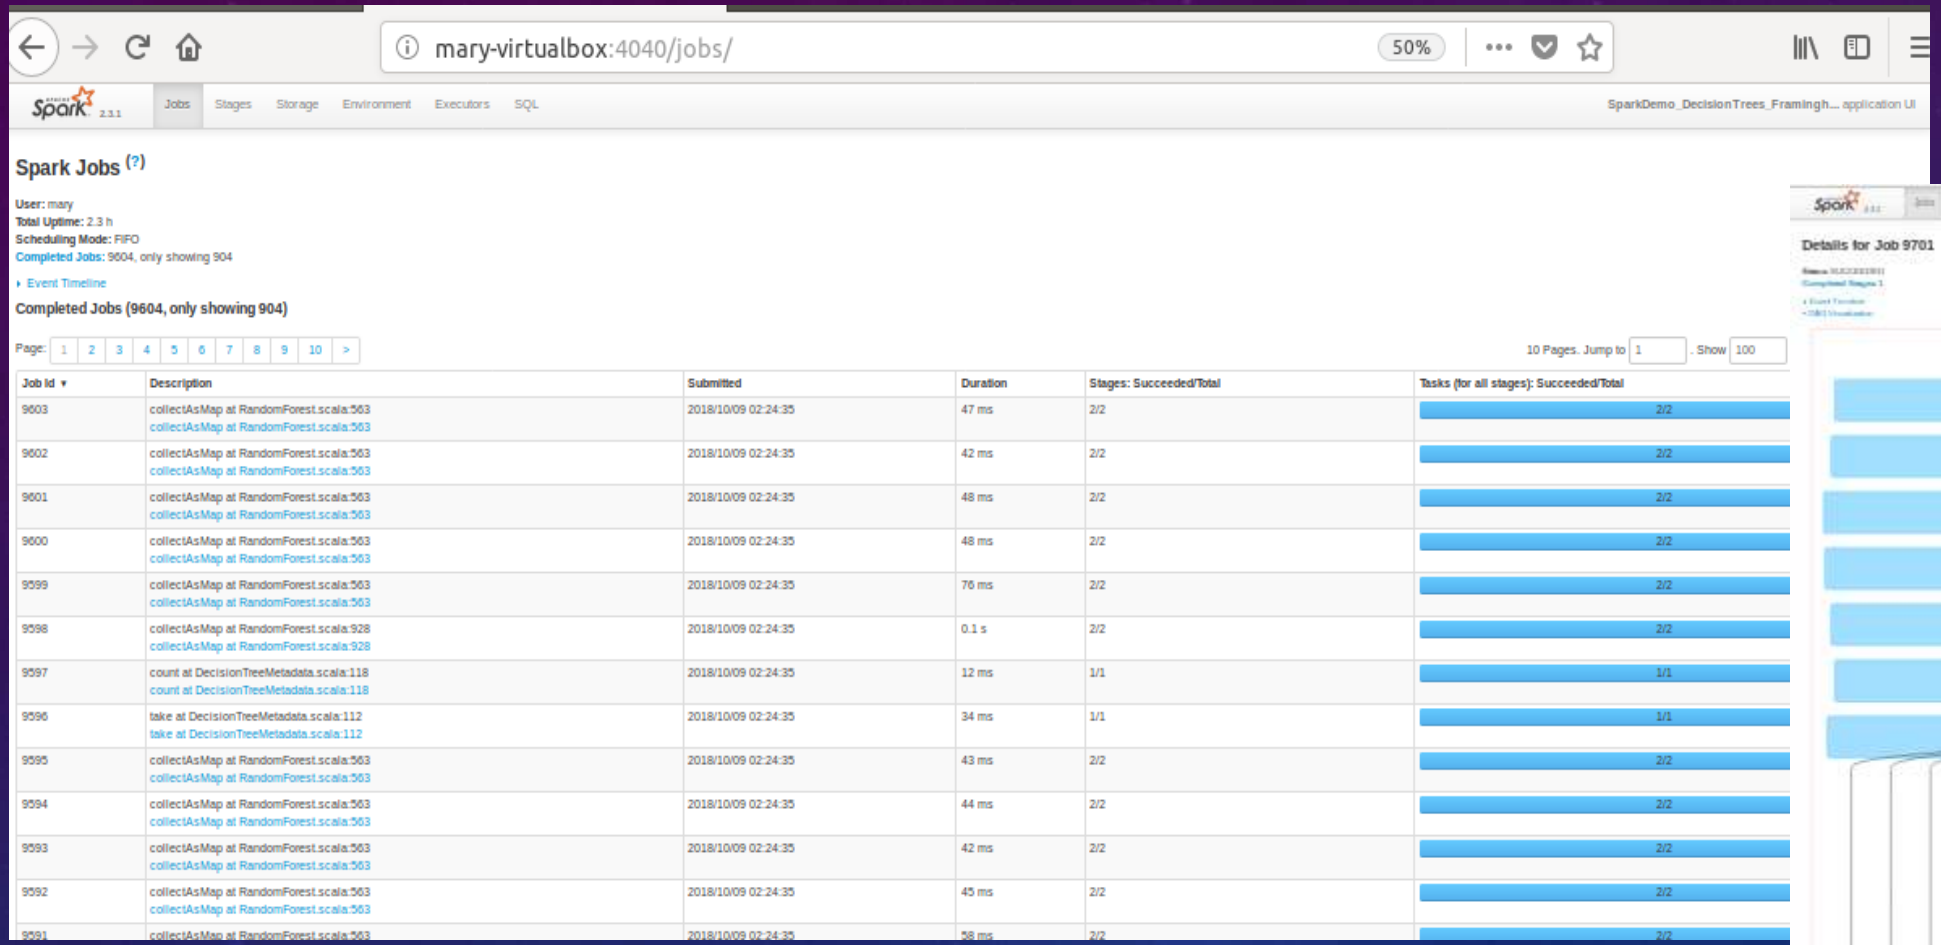

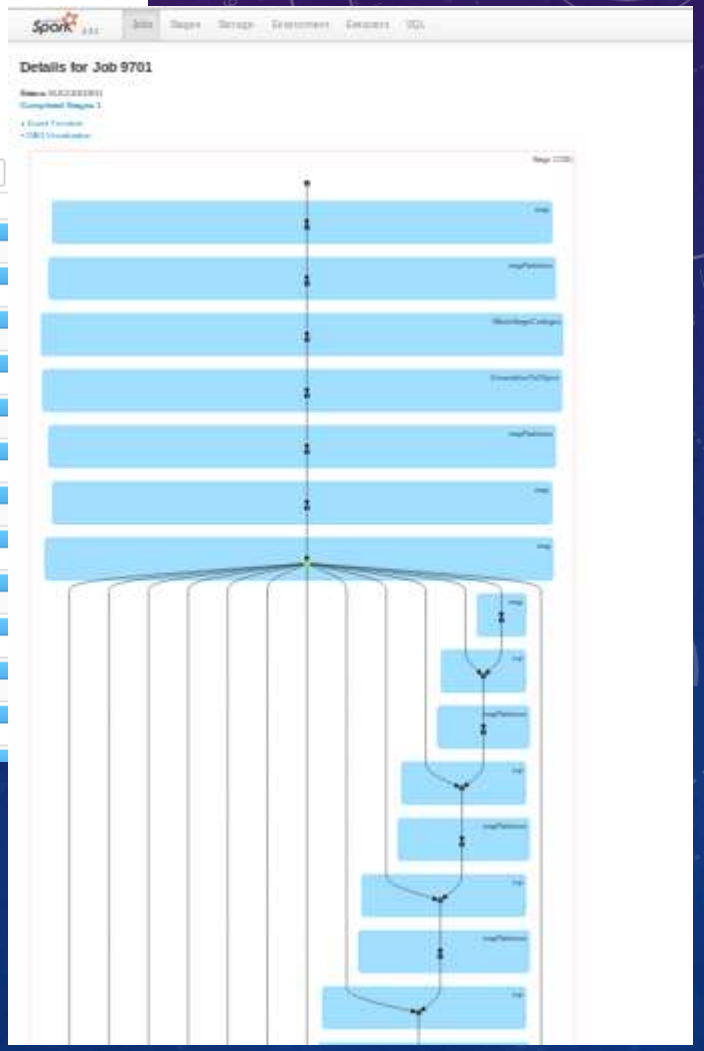

Deployment stage: Stopping Spark Session

Stop Spark Session

```
In [13]: 1 spark.stop()
```

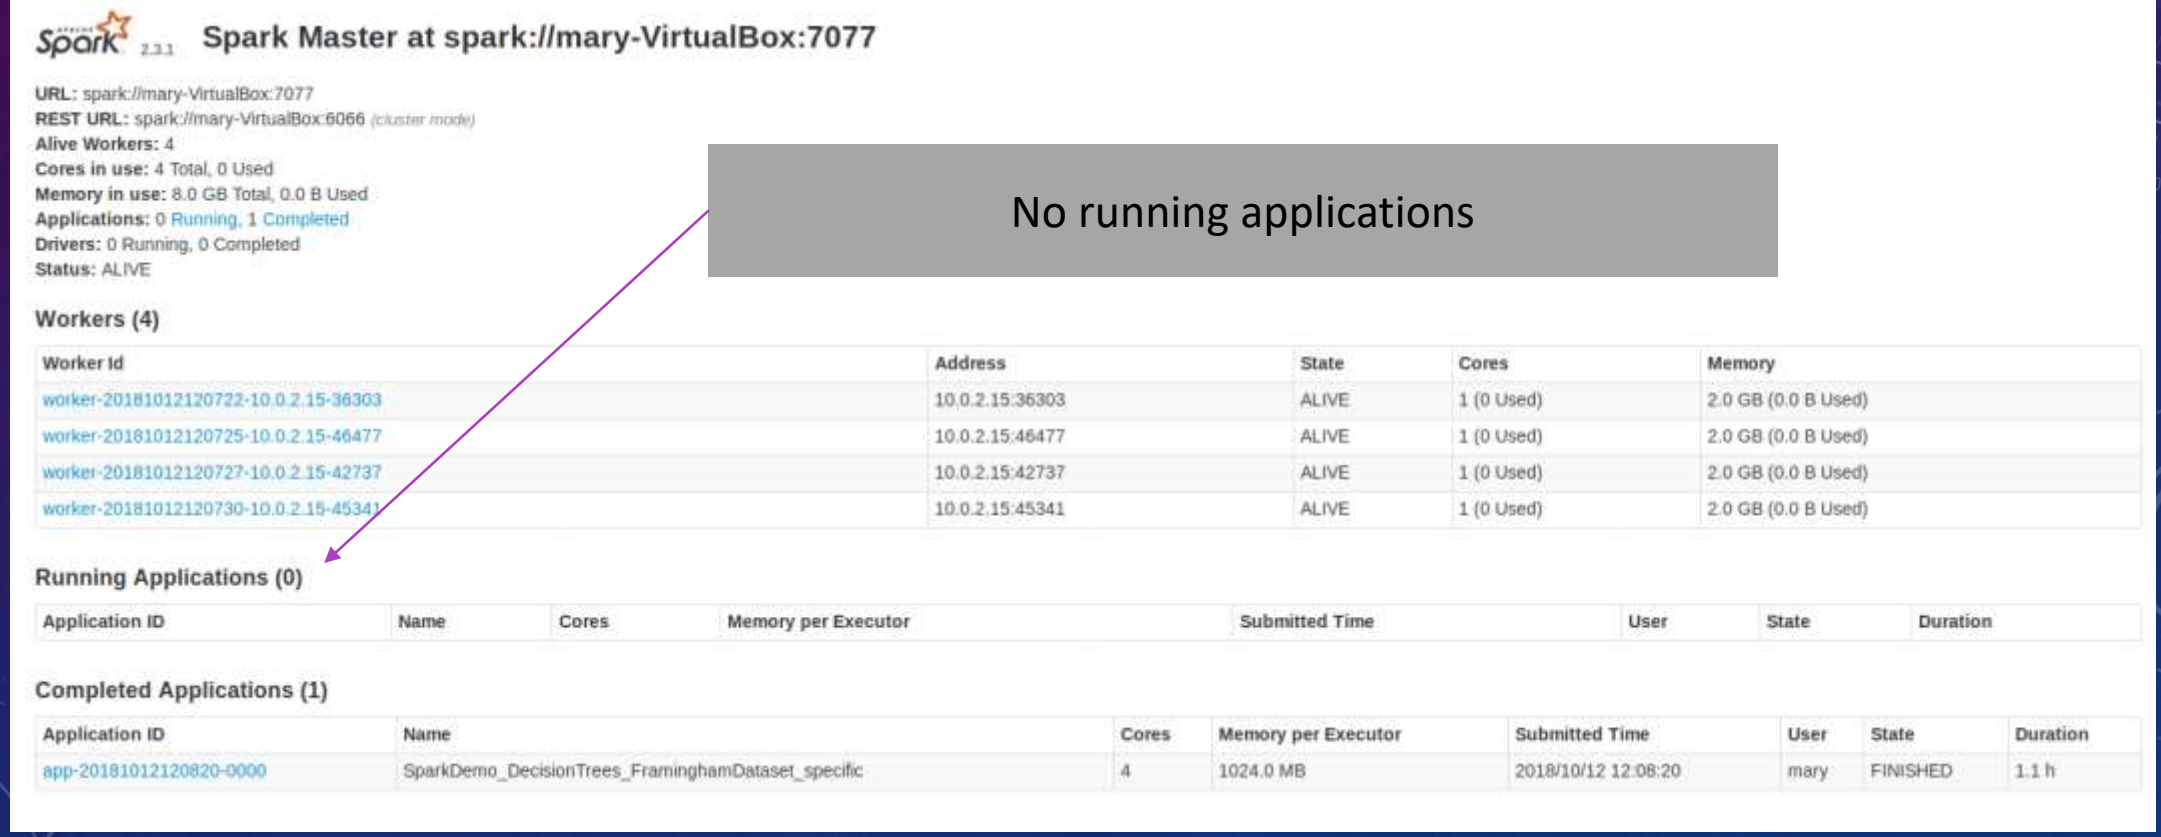

FINDINGS/OUTCOMES

Project Approach/Methodology

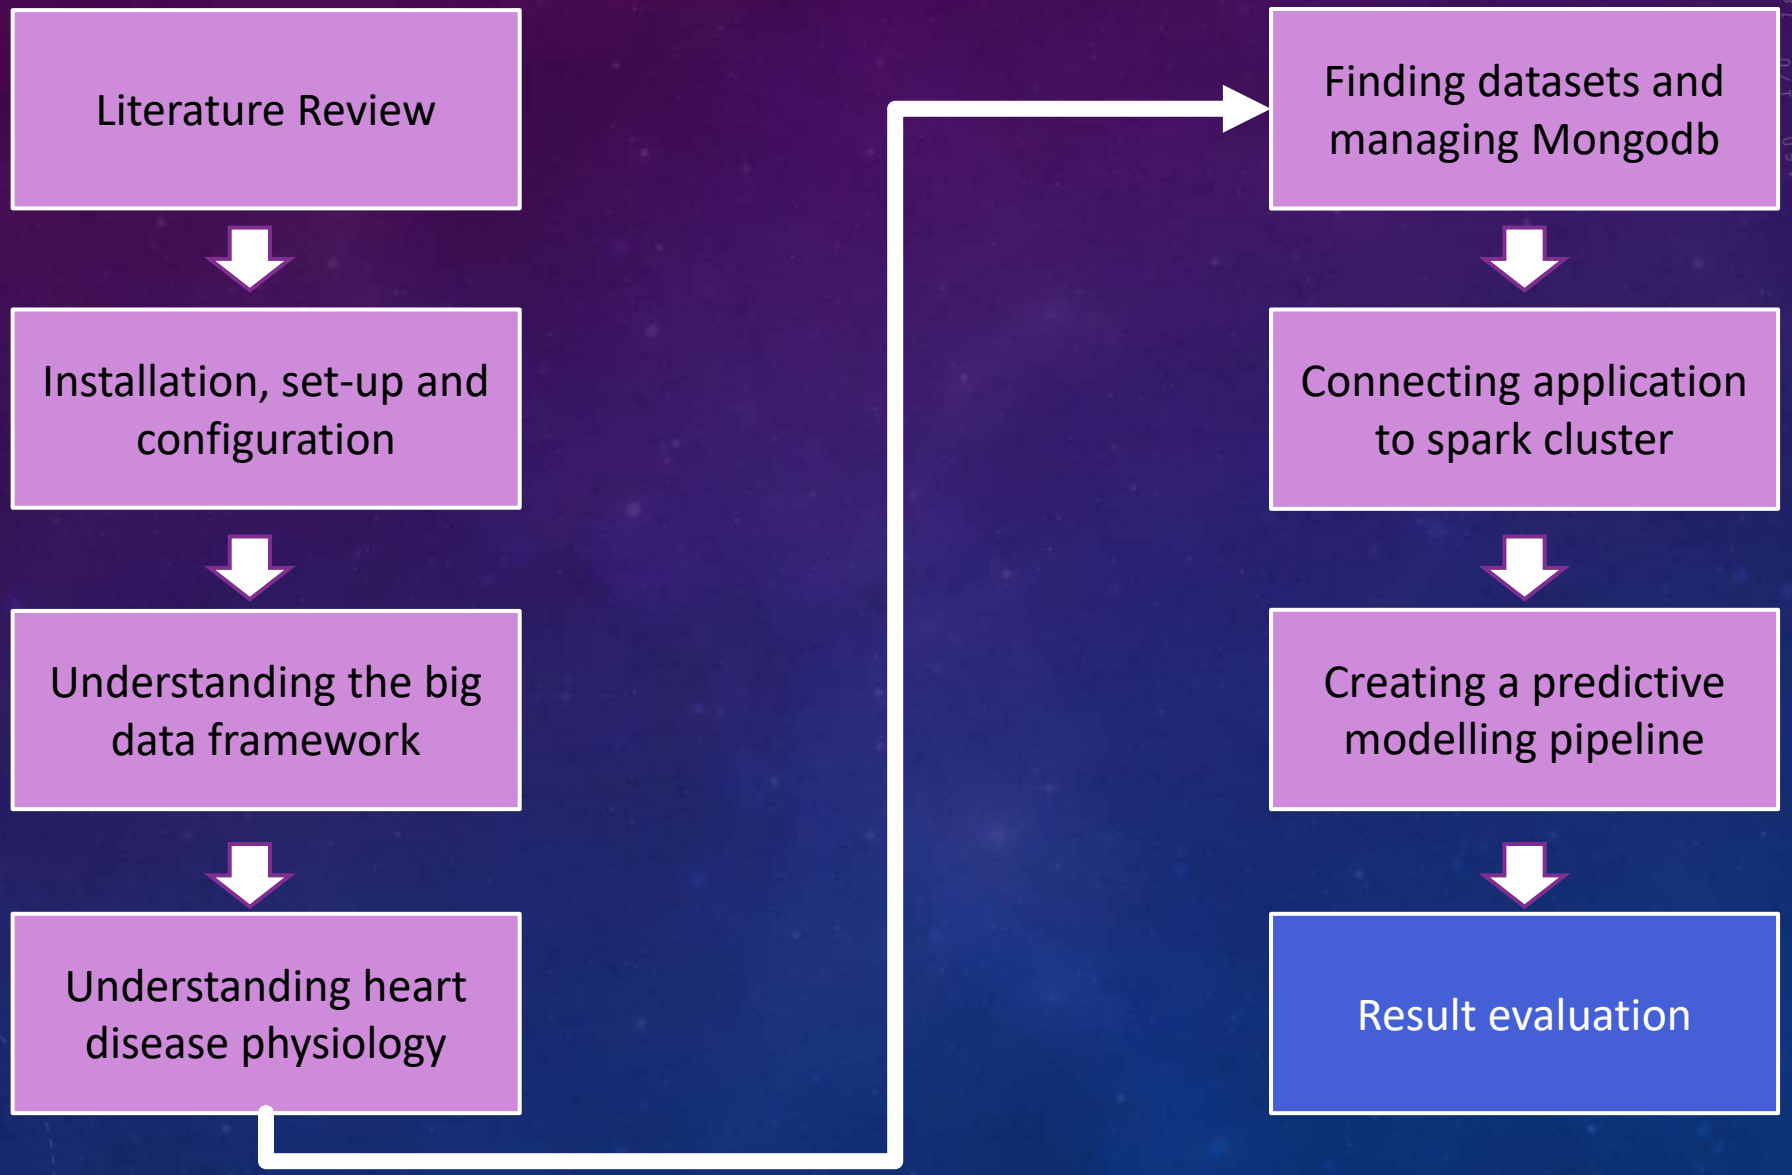

Findings/outcome

Setting Up DataFrame for Machine Learning

In [9]:
```
# remove educaiton and id columns
new_data=data.drop('education').drop('_id')
```

Convert independent variables (fields) to features

In [10]:
```
features=['sex','age','currentSmoker','cigsPerDay','BPMeds','prevalent
          'diabetes','totChol','sysBP','diaBP','BMI','heartRate',
assembler = VectorAssembler(inputCols=features,outputCol='features')
transformed_data = assembler.transform(new_data)
final=transformed_data.select('features','TenYearCHD')
```

Split data training 75%, test 25%

In [11]:
```
train_data,eval_data = final.randomSplit([0.75,0.25], seed=123)
```

Evaluation metrics

In [12]:
```
# evaluation metrics
acc_evaluator = MulticlassClassificationEvaluator(labelCol="TenYearCHD
prec_evaluator = MulticlassClassificationEvaluator(labelCol="TenYearCH
rec_evaluator = MulticlassClassificationEvaluator(labelCol="TenYearCHD
f1_evaluator = MulticlassClassificationEvaluator(labelCol="TenYearCHD"
```

In [301]:
```
# Decision Tree classifier
def DecisionTree():
    dtc = DecisionTreeClassifier(labelCol='TenYearCHD',featuresCol='fe
    dtc_model = dtc.fit(train_data)
    dtc_predictions = dtc_model.transform(eval_data)
    dtc_acc = acc_evaluator.evaluate(dtc_predictions)
    dtc_prec = prec_evaluator.evaluate(dtc_predictions)
    dtc_rec = rec_evaluator.evaluate(dtc_predictions)
    dtc_f1 = f1_evaluator.evaluate(dtc_predictions)
```

In [17]:
```
# GBT
def GBT():
    gbt = GBTClassifier(labelCol='TenYearCHD',featuresCol='features', maxIter=100, maxDepth=5, maxBins=32, \
                        minInstancesPerNode=1, cacheNodeIds=True)
    gbt_model = gbt.fit(train_data)

    # evaluate model
    gbt_predictions = gbt_model.transform(eval_data)
    gbt_acc = acc_evaluator.evaluate(gbt_predictions)
    gbt_prec = prec_evaluator.evaluate(gbt_predictions)
    gbt_rec = rec_evaluator.evaluate(gbt_predictions)
    gbt_f1 = f1_evaluator.evaluate(gbt_predictions)
    print('A ensemble using GBT accuracy  : {0:2.2f}%  precision:{1:2.2f}%  recall:{2:2.2f}%  f1:{3:2.2f}%'.format(gbt

    print("*" *80)
    total=gbt_predictions.count()
    POP=gbt_predictions.filter(gbt_predictions['prediction']==1).count()
    PON=gbt_predictions.filter(gbt_predictions['prediction']==0).count()
    CP=gbt_predictions.filter(gbt_predictions['TenYearCHD']==1).count()
    CN=gbt_predictions.filter(gbt_predictions['TenYearCHD']==0).count()
    TP=gbt_predictions.filter(gbt_predictions['TenYearCHD']==1).filter(gbt_predictions['prediction']==1).count()
    TN=gbt_predictions.filter(gbt_predictions['TenYearCHD']==0).filter(gbt_predictions['prediction']==0).count()
    FP=gbt_predictions.filter(gbt_predictions['TenYearCHD']==0).filter(gbt_predictions['prediction']==1).count()
    FN=gbt_predictions.filter(gbt_predictions['TenYearCHD']==1).filter(gbt_predictions['prediction']==0).count()
    Accuracy=(TP+TN)/total
    tpr=TP/CP
    fnr=FN/CP
    fpr=FP/CN
    tnr=TN/CN
    prev=CP/total
    ppv=TP/POP
    fdr=FP/POP
    fora=FN/PON
    npv=TN/PON

    print("total pop:", total)
    print("CP:",CP)
    print("CN:",CN)
    print("POP:", POP )
    print("PON:", PON)
    print("TP:",TP)
    print("TN:",TN)
    print("FP:", FP)
    print("FN:", FN)
    print("Accuracy:{0:2.2f}%".format(
    print("Recall/True positive rate:{0
    print("False negative rate:{0:2.2f}
    print("False positive rate:{0:2.2f}
    print("True negative rate:{0:2.2f}%
    print("Prevalence:{0:2.2f}%".format
    print("Precision/Positive predicted
```

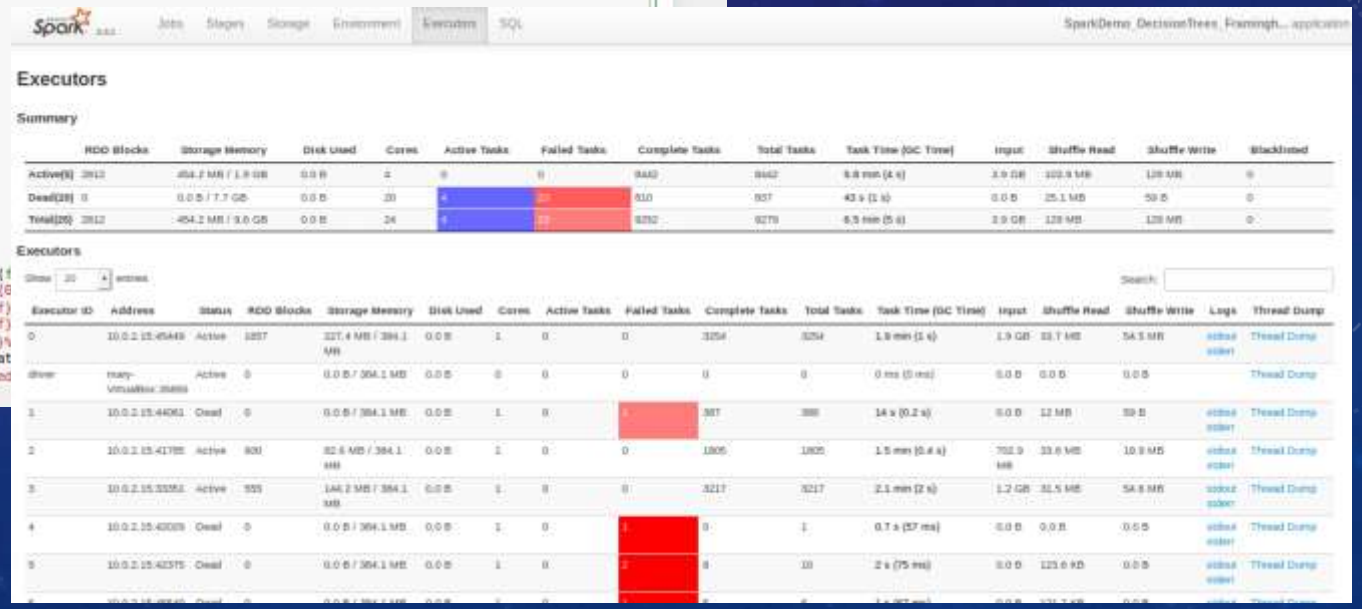

Findings/outcome

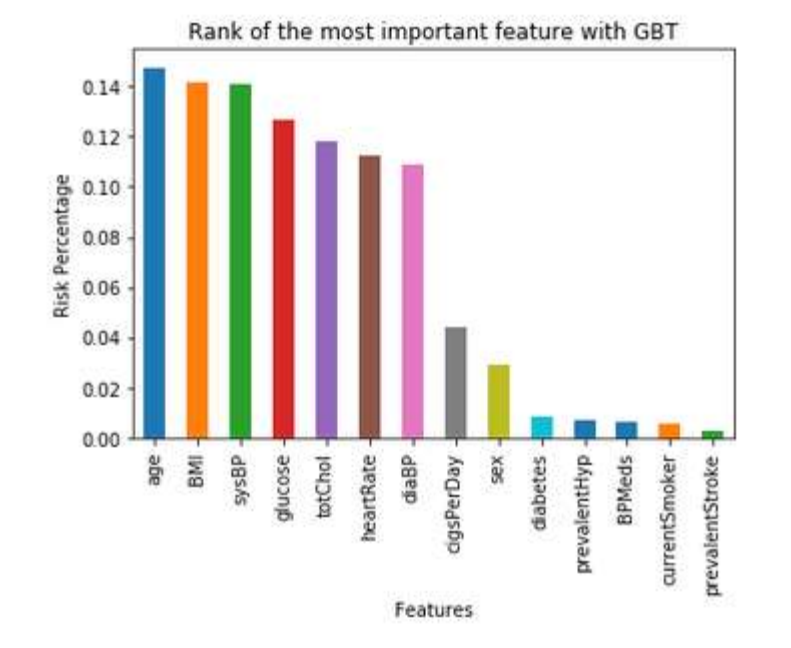

Accuracy	Precision	Recall	f1
85.08%	81.30%	85.08%	81.59%

Age, body mass index, blood sugar level, systolic blood pressure and cholesterol are the most important determining factor to assess risk of having coronary heart disease

Findings/outcome

		Ground Truth			
	Total Population 905	Condition Positive 135	Condition Negative 770	Prevalence 14.92%	
Prediction	Predictive Outcome Positive 135	True Positive 114	False Positive 21	Positive Predictive Value 84.44%	False Discovery Rate 15.5%
	Predictive Outcome Negative 770	False Negative 21	True Negative 749	False Omission Rate 2.72%	Negative Predictive Value 97.27%
Accuracy 85.08%		True Positive Rate 85.08%	False Positive Rate 2.73%		
		False Negative Rate 15%	True Negative Rate 97.27%		

Digest features and determine which specific subsets or range is likely contributing to coronary heart disease.

Findings/outcome: Age group

```
Here are the results!
A ensemble using GBT accuracy  : 86.17%  precision:74.26%  recall:86.17%  f1:79.77%
################################################################################
********************************************************************************
       risk features
0  0.379652      60s
1  0.244219      50s
2  0.208924      40s
3  0.166460      30s
4  0.000745      70s
################################################################################
```

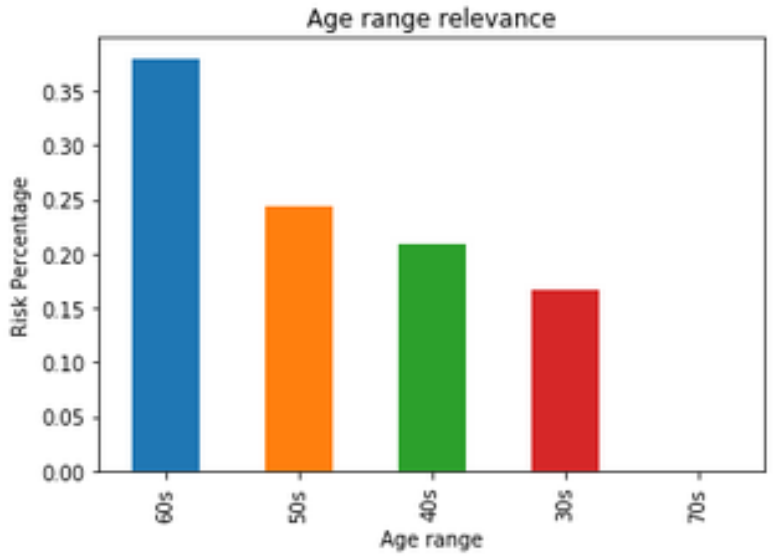

Findings/outcome: by Glucose level

```
Here are the results!
--------------------------------------------------------------------------------
A single decision tree accuracy: 85.41%  precision:73.16%  recall:85.41%  f1:78.82%
--------------------------------------------------------------------------------
A randomForestEnsemble accuracy: 85.55%  precision:73.18%  recall:85.55%  f1:78.88%
--------------------------------------------------------------------------------
A ensemble using GBT accuracy  : 85.41%  precision:73.16%  recall:85.41%  f1:78.82%
################################################################################
       risk features
0  0.255947  125-174
1  0.182892  100-124
2  0.139361    40-99
3  0.124695  275-324
4  0.113635  325-399
5  0.098622  225-274
6  0.084848  175-224
7  0.000000     400+
################################################################################
```

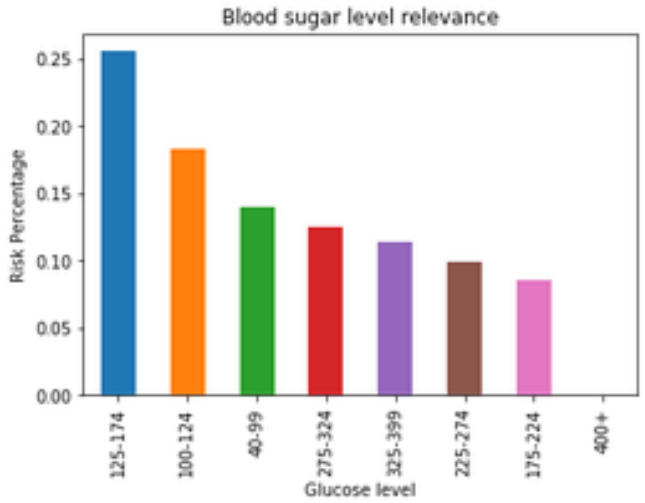

Findings/outcome: by Glucose level

```
A RandomForestEnsemble accuracy: 86.60%  precision:74.99%  recall:86.60%  f1:80.37%
--------------------------------------------------------------------------------
A ensemble using GBT accuracy  : 86.60%  precision:74.99%  recall:86.60%  f1:80.37%
################################################################################
       risk features
0  0.290246      1-9
1  0.210968    40-49
2  0.129764    30-39
3  0.124156    20-29
4  0.110447    50-59
5  0.066960    10-19
6  0.062431    60-69
7  0.005028      70+
################################################################################
```

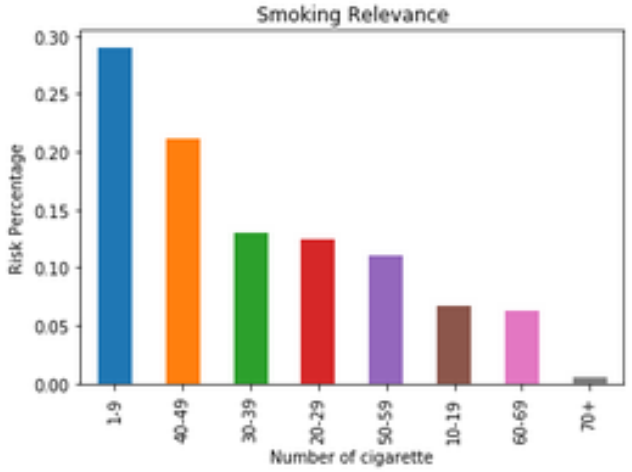

IMPLICATIONS

Implications

- Health care authorities may use the findings to focus their health education not only to elderly but also the young individuals as the risk for heart disease increased significantly higher as they grow older.
- Regulatory commissions to study sugar content of consumable goods and reduce use of sugar.
- Promote healthy lifestyle by quitting cigarette smoking as oppose to cutting down.
- Although the dataset is not a big data, the framework could work using the big data. Thus, applicable for scalable projects in predicting heart disease.

THANK YOU FOR LISTENING